

A RESEARCH ESSAY OF THE SERIES ON INTEGRATED AUTONOMOUS INTELLIGENCE

THE EXO-BRAIN

How AI Is Expanding Our Capacity
to Think, Build, and Act

EXO-AWARENESS · EXO-TOOLING · EXO-PRODUCTION

*“Artificial intelligence is moving beyond the production of answers
and entering the domains of construction, execution, memory,
and agency.”*

ALEJANDRO REYNOSO

JUDGE BUSINESS SCHOOL · UNIVERSITY OF CAMBRIDGE · JULY 2026

The Exo-Brain

Opening proposition

Artificial intelligence is moving beyond the production of answers and entering the domains of construction, execution, memory, and agency.

Dear reader

Dear reader, have you noticed what happens when you open your preferred AI platform?

You enter because you want to ask a question, develop an idea, improve a text, or understand something difficult. But around the familiar conversation window, new options have begun to appear. You see names such as Codex, Work, Claude Code, Cowork, or Antigravity. Perhaps an animated message invites you to try one of them. Perhaps you look at it briefly, decide that it is probably intended for programmers, and return to the reassuring simplicity of the chat.

You are not alone.

Most of us entered the age of generative AI through conversation. We learned that we could write a question in ordinary language and receive an intelligent response. We began using chatbots to explain concepts, summarize information, brainstorm possibilities, draft messages, analyze documents, and think through decisions. Chat became our window into artificial intelligence.

But something much larger is now happening around that window.

The major AI companies appear to be organizing the relationship between people and artificial intelligence into three increasingly distinct modalities: Chat, Code, and Work. Anthropic offers Claude, Claude Code, and Claude Cowork. OpenAI increasingly differentiates ChatGPT, Codex, and ChatGPT Work. Google's architecture is less symmetrical, but Gemini, Google Antigravity, and its expanding ecosystem of connected agents express a similar direction.

These are not merely new buttons or marketing categories. They represent three different relationships between human intention and machine intelligence:

Chat helps us think.
Code helps us build.
Work helps us accomplish.

Behind this apparently simple distinction lies a profound expansion of the boundaries of intelligence. Artificial intelligence is moving beyond the production of answers and entering the domains of construction, execution, memory, and agency.

To understand the significance of this transformation, it may be useful to consider the original instrument through which human intelligence entered the world: the human brain.


From awareness to civilization

The biological brain gave human beings the capacity to perceive their surroundings, remember experience, anticipate danger, recognize opportunities, and preserve themselves. But human intelligence did not stop at awareness.

Human beings used their intelligence to create tools. A stone could become a cutting instrument. A branch could become a lever. Materials found in nature could be transformed into extensions of the body and instruments of intention.

Those tools were then incorporated into organized activity. They were used to hunt, cultivate land, build shelters, produce food, transport materials, and eventually construct increasingly complex societies. Human beings did not merely understand their environment: they developed instruments through which they could transform it.

This was not a simple chronological sequence. Awareness, language, toolmaking, cooperation, and production co-evolved over immense periods. Nevertheless, the structure is recognizable:



We perceived the world. We constructed instruments. We used those instruments to produce outcomes.

Something structurally similar is now occurring in the evolution of our relationship with artificial intelligence. The difference is that we are adding a new cognitive system to the equation—an external layer of memory, reasoning, construction, and action.

We are beginning to build what might be called an exo-brain.

The meaning of an exo-brain

The idea is inspired by the exoskeleton.

A robotic exoskeleton does not replace the human body. It surrounds the body, detects its intentions, supports its movements, and amplifies its physical capabilities. The person remains the source of purpose and direction, while the external system provides additional strength, reach, endurance, or precision.

An exo-brain would perform an analogous function for cognition.

It would not replace the human mind, nor would it necessarily possess consciousness of its own. It would couple human purpose, experience, values, and judgment with external memory, computational reasoning, environmental perception, tool creation, and bounded execution.

An exoskeleton expands what the body can do. An exo-brain expands what the person can perceive, remember, understand, build, and accomplish.

This gives us a more precise interpretation of the three emerging AI modalities:

Table 1. Human capacities and their emerging AI extensions.

Human capacity	AI extension	Current modality
Awareness and judgment	Exo-awareness	Chat
Toolmaking and invention	Exo-tooling	Code
Work and coordinated action	Exo-production	Work

These dimensions are not separate products placed side by side. Together, they form an expanding intelligence manifold: a space in which human and artificial capabilities can be connected in new combinations.

Chat and the expansion of awareness

The first dimension is the one we already know best.

When we use ChatGPT, Claude, or Gemini, we enter into a conversation with a system capable of interpreting language, examining alternatives, detecting relationships, and generating explanations. We ask it to clarify an idea, compare scenarios, challenge an assumption, or help us determine what something means.

The characteristic request of Chat is:

Think with me.

The most important output is therefore not text. It is cognitive mediation.

Chat can widen the field brought before our attention. It can compare sources, retrieve relevant concepts, identify contradictions, generate competing interpretations, and expose possibilities that we might not have considered independently.

This is the beginning of exo-awareness.

The term must be used carefully. Exo-awareness does not mean that the AI is necessarily self-aware. It means that the system expands the human user’s effective field of awareness. A telescope does not itself see, yet it extends human vision. A microscope does not understand biology, yet it reveals a world that biological vision cannot perceive unaided.

Similarly, an AI system can enlarge the informational and conceptual field available to human consciousness without possessing human consciousness itself.

Exo-awareness can help us notice:

- ▶ what we have forgotten;
- ▶ what has changed;
- ▶ what contradicts our assumptions;
- ▶ what patterns connect apparently separate subjects;

- ▶ what risks are emerging;
- ▶ what opportunities remain invisible;
- ▶ what questions we have failed to ask.

But ordinary Chat remains largely episodic. We have a valuable conversation, obtain an insight, and then move on. Days later, the insight may be buried among hundreds of previous chats. The model may have helped us think, but the result has not necessarily become part of an enduring intellectual architecture.

This is where the idea of the second brain becomes essential.

From accumulated knowledge to a second brain

Consider the intellectual material you have already accumulated.

You have read books and articles. You have created presentations, documents, spreadsheets, notebooks, and applications. You have conducted hundreds of AI conversations. You have developed arguments, revised opinions, made decisions, and discovered relationships among ideas.

Yet much of that intelligence remains fragmented.

Some of it is in your biological memory. Some lies in folders whose contents you no longer remember. Some is buried in email. Some remains trapped in old chats. Some exists in notebooks, cloud drives, websites, or handwritten notes. The knowledge exists, but it does not operate as an integrated system.

A second-brain architecture begins by externalizing and organizing that intellectual history.

An Obsidian vault is a compelling example. It can contain human-readable notes, sources, concepts, projects, decisions, and connections. Its structure reflects not merely what the world knows, but how a particular person has encountered and interpreted that knowledge.

Your Obsidian vault might contain the books you have read, the courses you have taught, the research you have conducted, the questions that remain unresolved, and the connections you perceive among subjects that other people might regard as unrelated.

The vault is therefore more than storage. It is an external representation of part of your intellectual identity.

But Obsidian alone is not yet an exo-brain. It is the memory substrate of one.

The broader architecture requires several connected layers:

- ▶ the human brain supplies purpose, experience, values, intuition, and judgment;
- ▶ the vault supplies persistent and structured memory;
- ▶ the language model supplies interpretation, synthesis, and reasoning;
- ▶ an orchestration layer controls retrieval, permissions, tools, and workflows;
- ▶ coding agents create new capabilities;
- ▶ Work agents use those capabilities to produce outcomes.

When these layers are connected, something qualitatively different emerges:

**The vault gives AI continuity.
AI gives the vault interpretive capacity.
Tools give the combined system the ability to build.
Agents give it the ability to act.**

This is the transition from a passive second brain to an active exo-brain.

Code and the expansion of toolmaking

The second dimension is construction.

Human beings have always transformed ideas into tools, but toolmaking has traditionally required specialized skills. In the digital world, programming became one of the most powerful forms of toolmaking. Code could create an application, automate a process, execute a financial strategy, control a machine, or connect entire organizations.

Yet programming languages also created a barrier. A person might understand a problem deeply but be unable to convert that understanding into software.

Coding agents are beginning to reduce that distance.

With Claude Code, Codex, or Google Antigravity, the human can describe an objective in natural language. The agent can inspect files, navigate a repository, modify code, execute commands, run tests, diagnose errors, and progressively construct the requested system.

The characteristic request of Code is:

Build this.

This is the deeper meaning of vibe coding. It is sometimes presented as programming without technical knowledge, but that interpretation is too superficial. At its best, vibe coding relocates the human contribution from syntax production to intention, architecture, evaluation, and direction.

The human specifies:

- ▶ what problem must be solved;
- ▶ what behaviour is expected;
- ▶ what constraints must be respected;
- ▶ what risks are unacceptable;
- ▶ what evidence will demonstrate success.

The agent translates those intentions into executable capability.

This is exo-tooling: using an external cognitive system not merely to operate a tool, but to participate in creating new tools.

Natural language becomes a tool for making tools.

The result may be a website, an algorithmic trading system, a research application, a personal dashboard, a knowledge-retrieval engine, or an agent connected to the user's second brain. Each new tool can then extend the capacity of the larger system.

Exo-tooling is therefore recursive. We use the exo-brain to create a tool that expands what the exo-brain itself can subsequently do.

Work and the expansion of productive agency

A tool, however, does not create an outcome merely because it exists.

The agricultural transformation required more than the invention of instruments. Tools had to be incorporated into systems of observation, planning, coordination, execution, and learning.

The same distinction applies to AI. Producing code is not the same as completing work. An application is a capability; it becomes productive only when it is applied to an objective.

This is the domain of Work.

The characteristic request is no longer “What does this mean?” or “Build this application.” It becomes:

Accomplish this objective.

A Work agent can gather context, inspect files, choose tools, divide a project into stages, execute multiple operations, evaluate intermediate results, and deliver a completed artifact for review.

The outcome may be a document, spreadsheet, presentation, dashboard, research report, or application. But it may also be a change in the surrounding environment: an organized repository, an updated dataset, a reconciled financial model, a monitored process, or a recurring workflow.

This is exo-production.

The term goes beyond productivity. Productivity usually means producing more with the same resources. Exo-production changes the scope of productive agency. It enables an individual to initiate and supervise projects that might previously have required several specialists.

The distinction is important:

Exo-productivity helps us perform existing work more efficiently.

Exo-production expands the range of work we can undertake at all.


From digital systems to physical action

Code and Work do not need to remain confined to digital reality.

Software can connect with sensors, cameras, scientific instruments, vehicles, drones, industrial machinery, and robots. Once this happens, the exo-brain acquires channels through which it can perceive and affect the physical world.

Its awareness may be extended through satellites or medical devices. Its actions may be executed through a robotic arm, an autonomous vehicle, or a physical exoskeleton.

The full system begins to resemble an extended perception–action loop:



The perception may occur through external sensors. The interpretation may involve an LLM. The memory may reside in an Obsidian vault. The action may be performed by software in the cloud or by a robot in the physical world.

The biological boundary of human agency has not disappeared, but it has expanded.

The recursive intelligence cycle

Exo-awareness, exo-tooling, and exo-production continuously reinforce one another.




Figure 1. The recursive intelligence cycle.

Expanded awareness reveals a need or opportunity. Exo-tooling constructs the capability required to address it. Exo-production applies that capability to the world. The resulting experience returns to persistent memory and improves future judgment.

This recursive cycle is the real source of transformation.

A chatbot answers. An exo-brain accumulates.

It remembers what was attempted, what succeeded, what failed, what the user preferred, and how the surrounding environment changed. Its value can therefore grow through use—not because it becomes an autonomous copy of the individual, but because the coupling between human intelligence and computational capability becomes richer.

One person, two coupled cognitive systems

It is tempting to say that we now possess two brains. The metaphor is valuable, but it requires precision.

We do not have two equivalent consciousnesses. We have one human agent increasingly operating through two coupled cognitive systems:

- ▶ an embodied biological brain;
- ▶ an external computational exo-brain.

The human contributes purpose, responsibility, lived experience, moral judgment, and the capacity to care about consequences. The exo-brain contributes persistent memory, computational reach, rapid synthesis, continuous monitoring, tool generation, and scalable execution.

The meaningful unit is increasingly the partnership.

This partnership is not automatically beneficial. Poor information can expand confusion rather than awareness. Carelessly generated code can multiply errors. Autonomous execution can amplify a mistaken objective. The greater the exo-brain's capacity to act, the more important human judgment becomes.

Governance therefore does not disappear. It becomes embedded in the architecture through permissions, verification, provenance, approval thresholds, stopping conditions, and accountability.

The goal is not maximum autonomy. It is appropriate agency under meaningful human direction.

Have you noticed what those buttons mean?

Let us return to the screen with which we began.

You open your AI platform and see Chat, Code, and Work. Perhaps until now they looked like additional features competing for your attention.

But they are better understood as doors into three dimensions of augmented intelligence:

- ▶ Chat asks you to expand what you can understand.
- ▶ Code asks you to expand what you can create.
- ▶ Work asks you to expand what you can accomplish.

Behind those doors lies a larger possibility. The many conversations you have conducted, books you have read, files you have created, and judgments you have formed need not remain scattered intellectual fragments. They can become the foundations of a persistent second brain.

Connected to AI reasoning, that second brain becomes a source of exo-awareness. Connected to coding agents, it becomes a source of exo-tooling. Connected to Work agents—and eventually to robots and physical systems—it becomes a source of exo-production and extended agency.

We are therefore not merely witnessing the arrival of better chatbots.

We are participating in the construction of a new cognitive architecture: one in which human intelligence can be externalized, preserved, interpreted, expanded, and translated into action.

The biological brain allowed human beings to become aware, invent tools, and build civilizations. The exo-brain may now expand each of those capacities. It can help us perceive more than we could perceive alone, construct what we could not previously build, and undertake work that once exceeded the reach of an individual.

The question is no longer simply, "What can this chatbot tell me?"

The more important questions are:

What have I already learned that my exo-brain could help me recover?

What could I build if language became my instrument of construction?

What could I accomplish if my knowledge, tools, and agents operated as one coherent system?

Those unfamiliar buttons are not noise around the chat window. They are signs that artificial intelligence is moving from conversation to construction, and from construction to action.

They are the entrance to the exo-brain.