

A RESEARCH MONOGRAPH OF THE SERIES ON INTEGRATED AUTONOMOUS INTELLIGENCE

THE QUANTUM BRAIN

Quantum Navigation, Deliberative Intelligence and
Governed Memory in LLM–Knowledge Graph Systems

QUANTUM COGNITION · NB17–NB21

*“The objective is not to observe every possible reasoning path. It is
to reshape the probability of discovering the paths that matter.”*

ALEJANDRO REYNOSO

JUDGE BUSINESS SCHOOL · UNIVERSITY OF CAMBRIDGE · JULY 2026

Abstract

A governed knowledge graph can preserve more institutional knowledge than any large language model can place in a single context window. The central difficulty therefore shifts from storage to navigation: which paths, communities, contradictions, precedents, and weak signals should become cognitively active for a particular question? Classical graph retrieval already offers semantic ranking, diffusion, beam search, typed traversal, community summaries, and LLM-directed exploration. Yet as path depth and branching increase, the number of admissible evidence combinations becomes combinatorial. This monograph asks whether quantum computation can contribute to that search space without relying on the misleading claim that a quantum computer “considers every answer at once.”

The proposed Quantum Brain separates four functions. A governed knowledge graph stores typed claims, decisions, outcomes, temporal validity, provenance, and authority. A semantic compiler converts a natural-language question into seeds, targets, admissible relations, context constraints, and an explicit epistemic objective. A quantum navigation engine then operates on the combinatorial substrate beneath language: quantum walks reshape amplitudes across a graph frontier; amplitude amplification can increase the probability of states recognized by an efficient oracle; and QAOA, annealing, or related optimization routines can select portfolios of paths under relevance, diversity, contradiction, provenance, risk, and token constraints. Finally, a deliberative society of LLM agents interprets the retained portfolio, while verification and human governance control any recursive write-back.

The monograph formalizes path superposition, measurement, context-portfolio optimization, hybrid deliberation, and governed state transition. It distinguishes quantum difference, quantum usefulness, and quantum computational advantage; treats data loading, state preparation, oracle construction, circuit depth, noise, sampling, classical optimization, and output reconstruction as first-order costs; and argues that the proper baseline is the strongest relevant classical pipeline rather than brute force. The distinctive research object is not a quantum LLM, but quantum-conditioned knowledge evolution: repeated use of a quantum or quantum-inspired path distribution may change which candidate relationships are proposed and therefore how institutional memory evolves.

Five Colab-compatible notebooks implement a transparent laboratory. NB17 builds a 38-node, 54-edge governed graph and enumerates 320 candidate paths. NB18 simulates continuous-time quantum walks. NB19 formulates context construction as a nine-variable QUBO and benchmarks a statevector QAOA prototype against exact, greedy, and simulated-annealing methods. NB20 converts selected paths into role-based deliberation. NB21 runs 48 recursive rounds with governance and a strict advantage-claim ladder. The validated prototype finds broader quantum-conditioned contexts and greater topological divergence, but lower end-to-end utility than the classical baseline under the declared objective. The negative result is central: the architecture is meaningful and testable, but present evidence does not justify a quantum advantage claim.

Keywords: quantum computing; large language models; knowledge graphs; quantum walks; QAOA; QUBO; retrieval-augmented generation; institutional memory; governance; ExoBrain; hybrid intelligence.

Central proposition

A scientifically defensible Quantum Brain is not a quantum language model that reads thousands of complete answers simultaneously. It is a hybrid cognitive architecture in which classical retrieval compiles a governed graph frontier, quantum or quantum-inspired algorithms evolve or optimize a distribution over admissible reasoning paths, LLM agents interpret a diverse evidence portfolio, and governance controls what may become institutional memory.

Series note. This monograph advances the ExoBrain research programme from induced knowledge structure and navigation-conditioned evolution to quantum-conditioned navigation. An Obsidian vault is treated as a practical instance of a governed, versioned knowledge graph. The theory is software-independent.

Research status. The accompanying notebooks NB17–NB21 run without an API key. Their simulator results demonstrate distributional difference and a complete benchmarking method; they do not establish quantum computational advantage.

Executive Orientation

The phrase *Quantum Brain* can denote either a serious computational hypothesis or an attractive category error. The category error imagines a quantum computer as a magical parallel mind: every path is explored, every answer is available, and a measurement simply reveals the best one. Quantum mechanics does permit coherent superpositions, but it does not grant free access to every amplitude. Measurement returns a sample. The useful work occurs before measurement, when a carefully designed unitary process creates interference that suppresses some outcomes and amplifies others.

That correction does not weaken the idea. It gives the idea an implementable form.

In the preceding Navigation-Conditioned Knowledge Evolution framework, a question does not reach the entire graph. A navigation policy constructs a context subgraph. The answer depends on that subgraph; validated write-back then changes the graph available to future questions. The new leap is to regard the space of admissible paths and evidence bundles as a computational object in its own right. If a graph has average branching factor b and useful reasoning depth d , the number of paths can scale roughly as b^d . Even after classical pruning, context construction may require a choice among a very large number of overlapping, contradictory, and differently sourced path combinations.

The Quantum Brain therefore assigns different tasks to different machines:

LLM : semantic compilation, interpretation, and synthesis, quantum engine : distributional search and combinatorial selection, knowledge graph : externalized, typed, provenance-bearing memory, governance : authority, validation, audit, and write-back control.

This is a hybrid accelerator architecture, not a replacement theory of intelligence. Its prospective strength is not merely speed. It is the possibility of producing a deliberately plural cognitive field: a portfolio containing strong support, strong contradiction, a causal chain, a historical analogue, a cross-community bridge, and a weak signal, all within a fixed token budget. A larger chatbot can still retrieve a narrow context. A Quantum Brain, if successful, would improve the *distribution over ways of reasoning*.

The monograph proceeds with disciplined asymmetry. It is ambitious about architecture and conservative

about advantage. Quantum walks have provided provable speedups for particular oracle-defined graph problems [Childs et al., 2003, Childs and Goldstone, 2004]; Grover search provides a quadratic query advantage when a suitable oracle exists [Grover, 1996]; and variational algorithms offer flexible hybrid formulations [Farhi et al., 2014, Cerezo et al., 2021]. None of these results automatically transfers to a classical institutional knowledge graph. Data loading may dominate. An oracle may be expensive to build. A shallow circuit may be too local. Noise and sampling may erase a theoretical gain. A strong classical heuristic may outperform the quantum routine. These are not footnotes. They define the empirical programme.

Contents

Abstract	ii
Executive Orientation	iii
1 The Leap Beyond Navigation	1
1.1 From stored graph to path space	1
1.2 Why the classical solution is not dismissed	1
1.3 A new object: quantum-conditioned evolution	2
2 What “Quantum” Means Here	3
2.1 Superposition is not a readable list	3
2.2 Three distinct meanings of parallelism	3
2.3 Amplitude amplification and the oracle problem	3
2.4 Quantum walks and graph topology	4
2.5 Variational optimization and QAOA	4
3 Architecture of the Hybrid Super Brain	6
3.1 The six-layer system	6
3.2 Layer 1: the governed knowledge substrate	6
3.3 Layer 2: the semantic compiler	7
3.4 Layer 3: the classical frontier	7
3.5 Layer 4: the quantum engine	7
3.6 Layers 5 and 6: deliberation and authority	8
4 The Mathematics of Path and Portfolio Search	9
4.1 Path scoring as an epistemic objective	9
4.2 Why the highest-scoring paths are not enough	9
4.3 Path encoding choices	10
4.4 Measurement and retained context	10
4.5 A portfolio rather than a winner	10
5 From Quantum Samples to Deliberative Intelligence	11
5.1 The quantum engine does not speak	11
5.2 A society of constrained interpreters	11
5.3 Synthesis as an explicit contract	12
5.4 The superpower is cognitive plurality	12
6 Governance of Quantum-Conditioned Memory	13
6.1 Authority remains outside the model and the circuit	13
6.2 The Quantum Navigation Record	13
6.3 A three-level claim ladder	14
6.4 Write-back gates	14

6.5	Recursive risk	14
7	When Could Quantum Advantage Be Real?	15
7.1	The access-model problem	15
7.2	The output problem	15
7.3	The strongest-classical-baseline rule	15
7.4	Possible favorable regimes	16
7.5	Hardware realism	16
8	The Five-Notebook Experimental Programme	17
8.1	Identification strategy	17
8.2	NB17: architecture and path encoding	17
8.3	NB18: quantum walks	17
8.4	NB19: QAOA and QUBO context portfolios	18
8.5	NB20: quantum-conditioned multi-agent deliberation	18
8.6	NB21: governed recursion and advantage benchmarking	18
9	Validated Prototype Results	19
9.1	Baseline path space	19
9.2	Quantum-walk distribution	19
9.3	QAOA portfolio benchmark	19
9.4	Deliberation	20
9.5	Recursive graph evolution	20
9.6	The negative utility result	20
10	Scientific Propositions and Falsifiable Hypotheses	21
10.1	Distributional hypotheses	21
10.2	Recursive hypotheses	21
10.3	Advantage hypotheses	21
11	Research Risks and Design Failures	22
11.1	Quantum theater	22
11.2	Objective laundering	22
11.3	Frontier censorship	22
11.4	Noise-induced novelty	22
11.5	Recursive self-justification	22
11.6	Security and confidentiality	23
12	Implementation Roadmap	24
12.1	Stage 0: quantum-ready classical architecture	24
12.2	Stage 1: quantum-inspired laboratory	24
12.3	Stage 2: statevector and noisy simulation	24
12.4	Stage 3: reduced physical experiments	24
12.5	Stage 4: fault-tolerant readiness	24
12.6	Stage 5: governed institutional deployment	25
13	Institutional and Epistemological Implications	26
13.1	The unit of cognition becomes a distribution	26
13.2	A computational separation of inquiry and authority	26

13.3	Memory as an endogenous topology	26
13.4	The value of disagreement	26
14	Conclusion	27
A	Compact Mathematical Architecture	28
A.1	State sequence	28
A.2	Minimum comparative estimand	28
B	Notebook-to-Theory Map	29
C	Minimum Audit Bundle	30
	Bibliography	31

The Leap Beyond Navigation

1.1 From stored graph to path space

The construction of a rich knowledge graph is a major institutional achievement. Documents that were once isolated become connected to entities, transactions, decisions, assumptions, shocks, outcomes, and contradictions. In an Obsidian architecture, those relationships can remain human-readable while also being machine-parsable. The graph gives the organization an externalized memory whose structure is more expressive than a flat archive.

Yet a graph is not an answer. For a given query q_t , the complete graph $\mathcal{G}_t$ generally exceeds the available context budget. A retrieval system must produce a smaller object:

$$C_t = \mathcal{N}(\mathcal{G}_t, q_t; \pi_t, B_t),$$

where π_t is the navigation policy and B_t is the retrieval budget. The preceding monograph established that π_t is epistemically consequential. Semantic retrieval, local expansion, diffusion, community search, typed traversal, contrastive search, and exploration expose different subgraphs to the same model. GraphRAG and graph-guided reasoning systems make this practical distinction increasingly visible [Edge et al., 2024, Sun et al., 2023, He et al., 2024, Gutiérrez et al., 2024].

The new problem begins inside $\mathcal{N}$. A policy may need to compare a large set of candidate paths:

$$\mathcal{P}(\mathcal{G}_t, q_t; d) = \{p = (v_0, e_1, v_1, \dots, e_\ell, v_\ell) : \ell \leq d, p \text{ admissible for } q_t\}.$$

Path admissibility can encode edge types, dates, authority classes, source restrictions, risk rules, and maximum depth. The number of paths grows quickly, and the context problem is not merely to choose a single highest-scoring path. It is to assemble a non-redundant portfolio whose elements play different epistemic roles.

1.2 Why the classical solution is not dismissed

Classical algorithms are extraordinarily capable. Exact enumeration solves small instances. Beam search, Monte Carlo tree search, graph diffusion, personalized PageRank, integer programming, dynamic programming, submodular maximization, simulated annealing, and learning-based retrieval can solve or approximate much larger ones. A claim that “there are thousands of paths” does not by itself motivate quantum computation; ordinary machines can inspect thousands or millions of candidates.

Quantum interest begins only when three conditions align:

1. the admissible path or portfolio space is genuinely combinatorial;

2. the objective and constraints can be encoded without eliminating the prospective gain;
3. a quantum routine offers a distributional or scaling property that remains useful after end-to-end costs are counted.

The appropriate research question is therefore not “Can quantum computing search many paths?” It is:

Research question

Can a quantum or quantum-inspired search process produce, under equal evidence and context budgets, a more useful distribution over nonlocal, diverse, contradictory, and decision-relevant reasoning paths than the strongest relevant classical pipeline?

1.3 A new object: quantum-conditioned evolution

Suppose a quantum navigation backend changes which paths are measured or retained. The LLM then sees a different evidence portfolio, proposes different candidate relationships, and submits a different write-back package to governance. If some candidates are approved, the future graph changes.

Let κ_t denote the quantum algorithm, encoding, circuit parameters, noise model, and sampling procedure. The context becomes

$$C_t^Q = \mathcal{M} \left[U_{\kappa_t} (\text{Encode}(\mathcal{G}_t, q_t, \Omega_t)) \right],$$

where Ω_t is the epistemic objective and $\mathcal{M}$ includes repeated measurement and classical post-processing. The governed transition is

$$\mathcal{G}_{t+1} = \mathcal{G}_t \oplus \mathcal{H} \left[\mathcal{W} \left(\mathcal{L}(q_t, C_t^Q), C_t^Q, \kappa_t \right) \right].$$

The term *quantum-conditioned knowledge evolution* names the dependence of the future authoritative graph on the quantum-conditioned path distribution, without implying that every component or every advantage is quantum.

Proposition 1.1 (Backend non-neutrality). *If two retrieval backends induce different probability distributions over admissible evidence portfolios, and the write-back proposal is a non-constant function of the retained portfolio, then the expected future graph differs across backends even when the initial graph, question, model, and governance rule are identical.*

The proposition is elementary, but institutionally important. A quantum engine would not merely accelerate an internal computation. Repeated use could become a force shaping institutional memory.

What “Quantum” Means Here

2.1 Superposition is not a readable list

For a finite admissible path set $\mathcal{P}$, an idealized encoding may create

$$|\psi_0\rangle = \sum_{p \in \mathcal{P}} \alpha_p |p\rangle, \quad \sum_{p \in \mathcal{P}} |\alpha_p|^2 = 1.$$

This expression is often described informally as representing many paths “at once.” The description is acceptable only if its operational limit is kept explicit. A measurement does not print the vector $(\alpha_p)_p$. It returns a basis outcome p with probability $|\alpha_p|^2$. Recovering a useful distribution requires repeated shots, and estimating small probabilities can be expensive.

The algorithmic opportunity lies in transforming the amplitudes:

$$|\psi^\star\rangle = U_T U_{T-1} \cdots U_1 |\psi_0\rangle,$$

so that paths satisfying a useful property receive more probability mass. Interference is therefore a mechanism for *probability engineering*, not an oracle of omniscience.

2.2 Three distinct meanings of parallelism

Discussion of a Quantum Brain should distinguish:

Representational parallelism. A quantum state can encode amplitudes over a large basis.

Dynamical parallelism. A unitary can transform the full state according to quantum dynamics.

Extractable classical information. Measurement yields limited classical outcomes whose usefulness depends on the algorithm and number of shots.

Confusing the first two with the third produces the misleading claim that all reasoning paths have been “considered and read.” A rigorous architecture instead asks whether the measured samples have better task properties than samples produced classically.

2.3 Amplitude amplification and the oracle problem

Grover’s algorithm shows that an unstructured search over N items can locate a marked item with $O(\sqrt{N})$ oracle queries rather than $O(N)$ classical queries [Grover, 1996, Montanaro, 2016]. In a knowledge system, the seductive analogy is to mark “good reasoning paths.” But the oracle must recognize goodness coherently and efficiently.

If path quality depends on semantic relevance, source reliability, temporal validity, novelty, contradiction, and governance restrictions, oracle construction may be more difficult than path search. Grover Adaptive Search extends the logic to polynomial binary optimization but emphasizes the gate and ancillary-qubit costs of arithmetic and constraint recognition [Gilliam et al., 2021]. A Quantum Brain must therefore treat the oracle as a designed, audited component, not as a free assumption.

2.4 Quantum walks and graph topology

Quantum walks are a natural candidate because the problem is topological. Continuous-time evolution can be written

$$|\psi(t)\rangle = e^{-iHt}|\psi(0)\rangle,$$

where H is derived from the adjacency matrix, Laplacian, or a marked-node Hamiltonian. Amplitudes propagate along graph structure and interfere. Quantum walks have produced exponential separations for particular oracle problems and optimal spatial search on certain graph families [Childs et al., 2003, Childs and Goldstone, 2004]; quantum-walk algorithms also provide important polynomial improvements in other settings [Ambainis, 2007].

These results do not imply uniform superiority on arbitrary knowledge graphs. Performance depends on spectral structure, connectivity, marking, initial state, and the access model. A graph with severe hub dominance, irregular weights, or costly classical-to-quantum encoding may not preserve the theoretical advantage. The research task is to identify which knowledge-graph motifs, if any, create a useful regime.

2.5 Variational optimization and QAOA

QAOA alternates problem and mixer unitaries:

$$|\gamma, \beta\rangle = \prod_{\ell=1}^p e^{-i\beta_{\ell}H_M} e^{-i\gamma_{\ell}H_C} |+\rangle^{\otimes n}.$$

A classical optimizer selects the parameters to improve the expected objective [Farhi et al., 2014]. QAOA is appealing for context portfolios because relevance, redundancy, diversity, budget, and policy constraints can be expressed as binary optimization terms. QUBO and Ising formulations are well developed [Lucas, 2014]. Warm starts and risk-sensitive objectives such as conditional value at risk may improve some instances [Egger et al., 2021, Barkoutsos et al., 2020].

Yet variational algorithms introduce their own costs: parameter optimization, shot noise, barren plateaus, circuit depth, hardware connectivity, and repeated circuit execution [McClean et al., 2018, Cerezo et al., 2021]. A shallow circuit can be strongly local. A favorable small statevector result may disappear on hardware or at scale.

Scientific boundary

The monograph uses quantum algorithms as candidate computational mechanisms. It does not infer that quantum mechanics supplies consciousness, understanding, judgment, or institutional authority. The term *brain* refers to the functional architecture of external memory, search, deliberation,

learning, and governance.

Architecture of the Hybrid Super Brain

3.1 The six-layer system




Figure 3.1. The Quantum Brain as a hybrid, recursive, governed architecture.

The architecture is intentionally asymmetric. Language remains classical. The graph remains authoritative outside the quantum processor. The quantum component receives a compressed computational representation. The output returns to classical systems for interpretation, verification, and governance.

3.2 Layer 1: the governed knowledge substrate

The vault should distinguish at least:

- ▶ source documents and extracted claims;
- ▶ entities and typed relations;
- ▶ decisions, alternatives, assumptions, and owners;
- ▶ outcomes and retrospective evaluation;

- ▶ contradictions and unresolved disputes;
- ▶ temporal validity and supersession;
- ▶ confidence, source reliability, and authority status;
- ▶ navigation and model-generation records.

The quantum engine does not repair weak knowledge architecture. If contradictory claims have already been collapsed, if provenance is missing, or if candidate knowledge has been mixed with authoritative knowledge, a more powerful search process can amplify the defect.

3.3 Layer 2: the semantic compiler

The LLM translates a question into a structured query specification:

$$\Xi(q) = (S, T, \mathcal{R}^+, \mathcal{R}^-, d, B, \tau, \rho, \Omega),$$

where S contains seed nodes, T contains targets, $\mathcal{R}^+$ and $\mathcal{R}^-$ define permitted and prohibited relations, d is maximum depth, B is the context budget, τ is the temporal rule, ρ is the risk tier, and Ω is the objective.

This compilation step is itself governable. A high-risk credit decision might require causal, contradictory, historical-outcome, and regulatory paths. A low-risk exploratory research question might permit greater novelty and longer cross-community paths.

3.4 Layer 3: the classical frontier

Encoding the complete organizational vault directly into quantum memory is unlikely to be efficient in the near term. The system should first construct a smaller frontier $\mathcal{G}_t^F \subset \mathcal{G}_t$ using classical retrieval and governance filters. This is not a concession that makes the quantum layer irrelevant. Hybrid computing routinely assigns each processor the work it performs best. It is, however, a source of selection bias: a quantum engine cannot discover a node excluded from its frontier.

The frontier record must therefore include the exclusion rule, graph hash, node and edge counts, and coverage diagnostics. Comparative experiments should vary frontier size to test whether any apparent quantum benefit is merely a by-product of classical preprocessing.

3.5 Layer 4: the quantum engine

The quantum engine should expose interchangeable backends:

Backend	Cognitive function	Primary scientific question
Quantum walk	Redistribute probability across topology; reach nonlocal or bridge regions	Does topology-sensitive interference improve valuable discovery?
Amplitude amplification	Increase probability of oracle-recognized states	Can “good path” recognition be implemented efficiently?
QAOA	Optimize a constrained binary context portfolio	Does the sampled portfolio improve the quality-budget frontier?
Quantum annealing	Search an Ising/QUBO energy landscape	Is the embedding and annealing schedule competitive end to end?
Amplitude estimation	Estimate expectations or scenario probabilities	Does quadratic sampling improvement survive state preparation?
Quantum-inspired	Reproduce structural ideas with tensor, stochastic, or spectral methods	Which gains arise from the idea rather than the hardware?

Table 3.1. Candidate computational roles for the quantum navigation layer.

3.6 Layers 5 and 6: deliberation and authority

A measured path is not a judgment. The retained portfolio is assigned to roles with different obligations. The advocate constructs the strongest supported case. The skeptic seeks invalidating evidence. The causal analyst tests direction and mechanism. The historian retrieves comparable decisions and outcomes. The risk officer emphasizes tail exposure, constraints, and reversibility. The weak-signal explorer examines peripheral or cross-community paths.

Their outputs are evidence-bound: every statement must reference retained path identifiers. A synthesizer can then produce a conclusion, strongest alternative, unresolved contradictions, confidence, sensitivity, and required additional evidence. The result remains analysis. Governance decides whether any proposed relation is approved for authoritative memory.

The Mathematics of Path and Portfolio Search

4.1 Path scoring as an epistemic objective

Let each admissible path p have features

$$\phi(p) = (r_p, e_p, c_p, n_p, v_p, h_p, u_p, \ell_p),$$

representing query relevance, evidence quality, causal coherence, novelty, contradiction value, historical value, uncertainty, and length. A transparent path score is

$$s(p) = \beta_r r_p + \beta_e e_p + \beta_c c_p + \beta_n n_p + \beta_v v_p + \beta_h h_p - \beta_u u_p - \beta_\ell \ell_p.$$

The weights β are a policy object. They should vary by task and risk tier, be recorded in a registry, and be subject to sensitivity testing. A system that conceals the weights behind “AI relevance” is not governable.

4.2 Why the highest-scoring paths are not enough

Suppose the top ten paths all repeat the same central claim through minor variations. Their individual scores may be high while their joint information value is low. Let $x_i \in \{0, 1\}$ indicate whether path p_i is selected. A context portfolio can maximize

$$F(x) = \sum_i s_i x_i - \lambda \sum_{i < j} \text{sim}(p_i, p_j) x_i x_j \\ + \eta \text{ Balance}(x) + \nu \text{ Coverage}(x) - \rho \left(\sum_i c_i x_i - B \right)^2.$$

The redundancy penalty discourages repeated evidence. Balance rewards inclusion of both supporting and challenging paths. Coverage rewards communities, edge types, or evidence classes. The final term enforces the token or cost budget.

This is naturally expressible as QUBO when features and constraints are reduced to linear and quadratic terms:

$$\max_{x \in \{0,1\}^n} x^\top Q x + b^\top x + \text{constant}.$$

Binary reduction is not free. Higher-order conditions may require ancillary variables. Hard constraints may be weakened into penalties whose coefficients must be calibrated. The notebook laboratory therefore

retains an exact feasible benchmark for its small instance.

4.3 Path encoding choices

At least four encodings are possible:

1. **Path-index encoding:** each basis state denotes one pre-enumerated path.
2. **Edge-decision encoding:** qubits indicate whether an edge is selected.
3. **Node-portfolio encoding:** qubits indicate which evidence nodes enter context.
4. **Sequential-walk encoding:** the state evolves over graph vertices or directed edges.

Path indexing simplifies small prototypes but inherits classical enumeration cost. Edge encoding can represent exponentially many combinations but requires connectivity and path-validity constraints. Node portfolios align with context selection but may lose order and causal direction. Walk encoding is topologically natural but converts amplitudes over nodes into paths only through sampling or additional memory.

Proposition 4.1 (Encoding conservation). *A prospective quantum advantage in the search stage is not an end-to-end advantage if classical frontier construction, state preparation, constraint encoding, or output reconstruction has equal or worse asymptotic cost than the avoided classical search.*

4.4 Measurement and retained context

Let $P_\kappa(z) = |\langle z | \psi_\kappa^\star \rangle|^2$. After M shots, the system obtains samples $z_1, \dots, z_M$. Classical post-processing validates constraints, removes duplicates, computes evidence metrics, and retains k states:

$$C^Q = \text{Retain}_k \left(\text{Validate} \{z_m\}_{m=1}^M \right).$$

The retained context is therefore quantum-conditioned, not purely quantum. The complete method includes state preparation, circuit execution, measurement, validation, and classical retention. All must be benchmarked.

4.5 A portfolio rather than a winner

The decisive conceptual change is from argmax to portfolio:

$$\text{ordinary retrieval: } p^\star = \arg \max_p s(p),$$

$$\text{Quantum Brain: } C^\star = \arg \max_{C: \text{cost}(C) \leq B} \text{EpistemicValue}(C).$$

An executive decision rarely fails because the system omitted the sixth most similar sentence. It fails because context lacked a contradiction, a causal bridge, a precedent, a regulatory constraint, or a weak signal. The portfolio objective makes those absences measurable.

From Quantum Samples to Deliberative Intelligence

5.1 The quantum engine does not speak

Quantum computation manipulates encoded states. It does not explain a balance-sheet exposure, interpret a legal obligation, or write an investment-committee memorandum. A semantic layer must translate measured paths back into human-readable evidence packets.

For each retained path p , the packet should contain:

- ▶ ordered node and edge sequence;
- ▶ source excerpts or claim summaries;
- ▶ dates and temporal validity;
- ▶ reliability and authority status;
- ▶ contribution to the portfolio objective;
- ▶ contradictory or supporting orientation;
- ▶ quantum backend, parameters, and sampling frequency.

5.2 A society of constrained interpreters

The role architecture is not intended to anthropomorphize computation. It creates separation of duties. Each role receives an obligation and a bounded evidence set. No role is allowed to invent absent evidence.

Role	Required contribution	Failure prevented
Advocate	Strongest supported case and enabling conditions	Excessive conservatism
Skeptic	Strongest invalidating path and falsification test	Confirmation bias
Causal analyst	Mechanism, direction, mediators, and confounders	Associative overreach
Historian	Comparable decisions, assumptions, and realized outcomes	Organizational amnesia
Risk officer	Tail risk, constraints, liquidity, reversibility	Average-case myopia
Weak-signal explorer	Peripheral nodes and cross-community bridges	Hub dominance

Table 5.1. Deliberative roles applied to one quantum-conditioned evidence portfolio.

5.3 Synthesis as an explicit contract

The synthesizer must produce a structured object:

$$\mathcal{A} = (\text{recommendation, evidence chain, alternative, contradictions, confidence, sensitivity, next evidence}).$$

Confidence should not be a stylistic adjective. It should respond to evidence reliability, portfolio coverage, disagreement among roles, robustness across backends, and sensitivity to objective weights. If the recommendation changes when a small weight perturbation replaces one measured path, the system should report instability.

5.4 The superpower is cognitive plurality

A conventional retrieval pipeline often collapses uncertainty early: select top- k documents, summarize, answer. The proposed system postpones collapse. It deliberately preserves alternative explanations long enough for them to be compared. Quantum mechanics is not required for cognitive plurality; classical ensembles can provide it. The quantum hypothesis is that interference or combinatorial sampling may improve the distribution from which plurality is constructed.

Functional definition of the Super Brain

The Super Brain is a governed system that can generate, retain, compare, and learn from a portfolio of materially different reasoning paths without confusing plurality with truth or machine generation with authority.

Governance of Quantum-Conditioned Memory

6.1 Authority remains outside the model and the circuit

A quantum sample has no epistemic authority merely because it emerged from a sophisticated algorithm. It is a computational proposal. Likewise, an LLM interpretation is candidate analysis. Authority arises through institutional rules: source quality, accountable ownership, validation, approval, and version control.

The knowledge states should remain distinct:

source claim → retrieved evidence → generated interpretation
→ candidate relation → validated authoritative relation.

Collapsing these states would allow a probabilistic navigation process to rewrite the memory that later justifies its own conclusions.

6.2 The Quantum Navigation Record

Every run should create a canonical record:

```
object_type: quantum_navigation_record
query_id: Q067
graph_version: G_21
frontier_hash: "...
encoding: path_portfolio_qubo_v1
algorithm: qaoa
backend: simulator
qubits: 18
circuit_depth: 6
shots: 4096
objective_weights:
  relevance: 0.25
  reliability: 0.20
  contradiction: 0.15
  diversity: 0.20
  novelty: 0.10
  provenance: 0.10
classical_baseline: beam_plus_submodular
retained_path_ids: [...]
noise_model: null
```

```
claim_status: exploratory
human_validation: required
```

The record is not merely for reproducibility. It explains why a particular evidence set became cognitively active.

6.3 A three-level claim ladder

Quantum rhetoric should be controlled by a formal ladder:

Level 1: quantum difference. The quantum-conditioned output distribution differs from the declared classical baseline.

Level 2: quantum usefulness. Under equal budgets, that difference improves a task metric that matters.

Level 3: quantum computational advantage. The improvement persists under physical execution, scaling, data-loading, error, sampling, and end-to-end resource accounting against the strongest relevant classical method.

Level 1 is not Level 2. A distribution can be different and worse. Level 2 is not Level 3. A statevector simulator can find a useful distribution while consuming vastly more classical resources than the baseline.

6.4 Write-back gates

A candidate relation should be rejected, deferred, or approved using explicit criteria:

- ▶ minimum source reliability and evidence breadth;
- ▶ presence of challenge evidence for high-impact supporting claims;
- ▶ temporal validity and non-supersession;
- ▶ absence of duplicate or circular support;
- ▶ stability across reasonable objective weights;
- ▶ human approval for high-risk domains;
- ▶ preserved navigation, model, and source provenance.

Governance can also restrict the quantum policy. A credit decision may prohibit weak-source novelty from offsetting regulatory evidence. A research task may permit broader exploration but forbid automatic write-back.

6.5 Recursive risk

If a quantum-conditioned sampler repeatedly finds cross-community bridges, approved write-back may make the graph less modular. That could improve institutional integration. It could also create spurious shortcuts. If it repeatedly amplifies central nodes, it may accelerate hub dominance. If it overvalues contradiction, it may generate artificial controversy. Topology audits must therefore compare graph trajectories across backends.

Proposition 6.1 (Recursive amplification). *When write-back probability increases with retrieval exposure, a backend that systematically raises the exposure of a node class will, other things equal, increase that class's future degree and future retrieval probability.*

When Could Quantum Advantage Be Real?

7.1 The access-model problem

Many quantum speedups assume efficient coherent access to input data. An institutional graph is classical, heterogeneous, frequently updated, and rich in text. Preparing the quantum state may require work proportional to the input size. If state preparation must be repeated for every query, the theoretical search advantage may be absorbed before the algorithm begins.

The architecture therefore uses classical compilation and a compact frontier. A future fault-tolerant implementation might exploit specialized quantum memory or structured state preparation, but such capability must be measured, not assumed.

7.2 The output problem

Quantum algorithms are most attractive when the desired output is small: a decision, an estimate, a sample, or a compact witness. They are less attractive when the user needs the full classical description of an exponentially large state. The Quantum Brain is compatible with this constraint because it needs a small portfolio of paths, not the entire distribution. Nevertheless, the number of shots needed to obtain sufficiently diverse, stable, and rare paths may be large.

7.3 The strongest-classical-baseline rule

Brute force is rarely the correct comparison. A credible benchmark should include:

- ▶ semantic plus typed beam search;
- ▶ personalized PageRank or diffusion;
- ▶ submodular or integer-programming portfolio selection;
- ▶ simulated annealing and modern local search;
- ▶ learning-to-rank or policy-guided traversal;
- ▶ quantum-inspired tensor or sampling methods.

The history of quantum-inspired algorithms shows why this matters. Some claimed quantum machine-learning advantages narrowed after classical algorithms adopted comparable data-access assumptions [Tang, 2019]. More broadly, classical learning from data can reduce the practical advantage of a quantum model on some tasks [Huang et al., 2021]. A Quantum Brain should treat dequantization as a success of the research programme: if quantum ideas produce a better classical cognitive architecture, institutional intelligence has still improved.

7.4 Possible favorable regimes

Quantum methods may become interesting where:

1. the frontier is structured and compact enough for efficient state preparation;
2. the path space is combinatorial but the required output is a small diverse sample;
3. graph topology supports favorable quantum-walk transport;
4. constraints admit efficient coherent evaluation;
5. the value of rare nonlocal discoveries is high;
6. queries recur over a relatively stable encoded frontier;
7. the strongest classical methods face a measurable quality-cost frontier.

7.5 Hardware realism

NISQ devices remain limited by noise, depth, connectivity, and measurement cost [Preskill, 2018]. Industry roadmaps anticipate larger fault-tolerant systems later in the decade; IBM, for example, publicly targets its Starling system for 2029 [IBM Quantum, 2025, 2026]. A roadmap is not a guarantee, and availability does not establish application advantage. The appropriate near-term strategy is backend interchangeability: exact classical, quantum-inspired, statevector, noisy simulator, and carefully reduced hardware experiments.

The Five-Notebook Experimental Programme

8.1 Identification strategy

The laboratory freezes:

- ▶ one governed baseline graph;
- ▶ the same question set;
- ▶ the same path and context budgets;
- ▶ the same feature definitions;
- ▶ the same governance rule;
- ▶ a fixed random seed;
- ▶ complete manifests and hashes.

The intervention is the navigation or portfolio-selection backend. This design isolates backend effects on retained context, deliberation, and recursive graph evolution.

8.2 NB17: architecture and path encoding

NB17 builds a synthetic investment-committee vault with companies, shocks, decisions, evidence, weak signals, and outcomes. It enumerates simple paths under a depth limit and calculates relevance, reliability, novelty, causal, contradiction, support, bridge, polarity, and cost features. Exact top- k , greedy, beam, and stochastic classical selectors create baseline portfolios.

The notebook produces the graph, query file, path catalogue, classical benchmark table, figure, and manifest. It also documents how to replace the synthetic graph builder with an Obsidian Markdown and front-matter parser.

8.3 NB18: quantum walks

NB18 constructs a weighted graph Laplacian, adds marked-query terms, and simulates

$$|\psi(t)\rangle = e^{-iHt}|\psi(0)\rangle$$

through eigendecomposition. It compares the probability distribution with a classical restart random walk, searches a transparent time grid, measures node portfolios, and applies a dephasing mixture. The purpose is to study distributional and topological behavior, not to claim hardware speedup.

8.4 NB19: QAOA and QUBO context portfolios

NB19 selects nine candidate paths for one query and creates a four-path context budget. The objective rewards path value, support-challenge balance, and non-redundancy. It solves the feasible problem exactly, constructs a greedy portfolio, runs simulated annealing, and performs a dependency-free statevector QAOA simulation at depths $p = 1$ and $p = 2$.

The exact solution is essential. It makes the QAOA approximation ratio interpretable and exposes infeasible high-probability states.

8.5 NB20: quantum-conditioned multi-agent deliberation

NB20 maps selected paths to six roles and creates an auditable decision memorandum. The default text generator is deterministic so the notebook requires no API key. A production deployment may replace it with approved LLM calls while preserving the same path packets, role contracts, and audit structure.

8.6 NB21: governed recursion and advantage benchmarking

NB21 gives classical and quantum-conditioned systems identical graph copies and runs 48 question rounds. Both propose candidate relations. A governance gate checks reliability, evidence breadth, contradiction coverage, and duplication before simulated approval. The notebook compares context quality, coverage, approval rates, graph topology, noise sensitivity, runtime, and a declared composite utility. It writes an audit ZIP with file hashes and the three-level claim ladder.

Validated Prototype Results

9.1 Baseline path space

The executed laboratory produced a 38-node, 54-edge governed graph and 320 candidate paths across four decision queries. Each query reached the cap of 80 admissible paths. The result is intentionally small enough for inspection, yet large enough to show why individual path ranking and portfolio construction are different problems.

Object	Count	Object	Count
Graph nodes	38	Graph edges	54
Queries	4	Candidate paths	320
QUBO variables	9	QUBO states	512
Context budget	4 paths	Recursive rounds	48

Table 9.1. Scale of the validated synthetic laboratory.

9.2 Quantum-walk distribution

For the AlphaBank acquisition query, the marked continuous-time walk reached its highest evidence-node probability mass of approximately 0.448 at evolution time $t = 3.514$ on the declared grid. At that point, the Jensen-Shannon distance from the classical walk was approximately 0.347. The measured top-node portfolios also retrieved weak signals in cases where the classical top portfolio did not.

This demonstrates Level 1, quantum difference in simulation. It does not show that the difference is more accurate, more useful, faster, or hardware-achievable.

9.3 QAOA portfolio benchmark

The exact feasible objective for the nine-variable, four-path portfolio was 9.432. The statevector QAOA prototype produced a best feasible high-probability state with objective 9.374, an approximation ratio of 0.994. Simulated annealing reached the exact feasible optimum in the validation run; the greedy method reached 9.389.

Method	Objective	Approximation ratio	Distinct nodes
Exact feasible	9.432	1.000	9
Greedy	9.389	0.995	8
Simulated annealing	9.432	1.000	9
QAOA feasible mode	9.374	0.994	8

Table 9.2. Validated context-portfolio benchmark. Values are simulator results on a synthetic instance.

The QAOA result is technically successful but not superior. On this instance, a strong classical heuristic matched the optimum.

9.4 Deliberation

The exact, greedy, and QAOA portfolios all generated a cautious recommendation on the acquisition query. After adding an explicit balance reward, the retained contexts included supporting and challenging paths. The role record also exposed a deficiency: no historical-outcome path survived the selected portfolio, so the historian correctly returned insufficient evidence. This is precisely the kind of omission that a governance rule can convert into retrieval escalation.

9.5 Recursive graph evolution

Across 48 rounds, the classical selector achieved an approval rate of 0.083; the quantum-conditioned selector achieved 0.208. The quantum-conditioned portfolios covered an average of 11.896 distinct nodes, compared with 6.750 for the classical selector. The final simple graphs contained 59 and 56 edges respectively, with an edge-set distance of approximately 0.152.

These results demonstrate that a broader sampling policy can materially change the graph trajectory under an identical governance gate. They do not show that the additional approved candidates are ultimately truer. Synthetic governance acceptance is a process metric, not ground truth.

9.6 The negative utility result

At zero injected sampler noise, the declared composite utility averaged 0.793 for the quantum-conditioned selector and 0.906 for the classical selector. The quantum-conditioned method gained diversity but lost mean score and runtime efficiency. The claim ladder therefore recorded:

- ▶ quantum difference: demonstrated in simulation;
- ▶ quantum usefulness: not demonstrated under the declared objective;
- ▶ quantum computational advantage: not claimed.

Why the negative result matters

The prototype prevents the architecture from becoming a slogan. It shows that broader exploration can be genuinely different yet still inferior on a declared utility function. Future work must improve the quantum-conditioned objective, encoding, and backend or identify graph regimes where diversity has higher decision value. The benchmark has done its job: it has made disappointment informative.

Scientific Propositions and Falsifiable Hypotheses

10.1 Distributional hypotheses

Hypothesis 10.1 (Nonlocal discovery). *For graph families with sparse cross-community bridges, a properly marked quantum-walk backend will increase the probability of retrieving valid cross-community evidence under an equal node-visit budget.*

Hypothesis 10.2 (Diversity-quality frontier). *A constrained quantum portfolio sampler will produce higher contradiction and community coverage than semantic top-k retrieval at the same context budget, without reducing evidence precision below a predeclared floor.*

Hypothesis 10.3 (Topology dependence). *Any quantum-walk usefulness will vary systematically with graph spectral properties, degree concentration, modularity, and bottleneck structure rather than appearing uniformly across knowledge graphs.*

10.2 Recursive hypotheses

Hypothesis 10.4 (Evolutionary divergence). *Initially identical governed graphs exposed to classical and quantum-conditioned navigation will diverge measurably in degree concentration, modularity, relationship diversity, and weak-signal retention after repeated approved write-back.*

Hypothesis 10.5 (Governance moderation). *Challenge-evidence requirements and topology audits will reduce retrieval-induced preferential attachment without eliminating the discovery advantage of exploratory backends.*

10.3 Advantage hypotheses

Hypothesis 10.6 (End-to-end crossover). *There exists a frontier size or path-complexity regime in which a physical quantum backend achieves higher task utility per unit of total resource than the strongest declared classical baseline after accounting for loading, optimization, shots, mitigation, and reconstruction.*

This last hypothesis is the most demanding and the least supported by the present prototype. Its value lies in defining what future evidence would need to show.

Research Risks and Design Failures

11.1 Quantum theater

The first risk is branding a classical stochastic pipeline as quantum merely because it borrows amplitude language. Quantum-inspired algorithms are legitimate, but they must be labeled accurately. The notebook laboratory distinguishes continuous-time statevector simulation, QAOA statevector simulation, and an amplitude-like quantum-conditioned sampler.

11.2 Objective laundering

A system can appear successful if the evaluation metric rewards exactly what its objective was designed to produce. If the portfolio objective rewards diversity and the benchmark celebrates diversity, the test is circular. External decision metrics, held-out questions, blinded expert evaluation, and realized outcomes are required.

11.3 Frontier censorship

Classical preprocessing can silently determine the result. If the frontier builder excludes peripheral communities, no quantum walk will recover them. Frontier recall, excluded-node audits, and size sensitivity must accompany quantum metrics.

11.4 Noise-induced novelty

Noise can increase entropy and produce apparently diverse samples. Diversity created by random corruption is not useful exploration. The experiment must distinguish structured cross-community discovery from uniformization. The NB18 dephasing sweep illustrates this point: entropy rises as the distribution moves away from the ideal state.

11.5 Recursive self-justification

If machine-generated relations are written back too easily, future retrieval can cite earlier machine outputs as if they were independent evidence. Provenance graphs must preserve derivation chains and prohibit circular corroboration.

11.6 Security and confidentiality

A real institutional vault may contain material non-public information, personal data, and regulated records. Quantum cloud execution would create an additional processing boundary. Frontier extraction, encryption, access control, residency, provider logging, and export restrictions require explicit review. Synthetic and simulator experiments should precede any sensitive deployment.

Implementation Roadmap

12.1 Stage 0: quantum-ready classical architecture

The first production milestone requires no quantum hardware. Build:

- ▶ a governed typed graph;
- ▶ a query compiler;
- ▶ an auditable path catalogue;
- ▶ a context-portfolio objective;
- ▶ strong classical baselines;
- ▶ deliberative role contracts;
- ▶ a navigation record and write-back gate.

This stage already improves the ExoBrain because it replaces top- k retrieval with plural, policy-aware context construction.

12.2 Stage 1: quantum-inspired laboratory

Use spectral propagation, tensor methods, stochastic amplitude-like sampling, QUBO solvers, and annealing. Determine which benefits arise from portfolio design and topology rather than quantum hardware. This stage is an empirical control, not an inferior placeholder.

12.3 Stage 2: statevector and noisy simulation

Run exact small circuits to validate encodings and compare distributions with known optima. Add realistic gate, readout, and dephasing noise. Estimate shots, depth, qubits, optimization calls, and reconstruction cost.

12.4 Stage 3: reduced physical experiments

Select small, structurally diagnostic graph motifs. Execute on available quantum hardware. Compare matched simulations and hardware outcomes. Report mitigation cost and uncertainty. Avoid scaling claims.

12.5 Stage 4: fault-tolerant readiness

Design algorithms whose advantage depends on clear access and output assumptions. Prepare logical-qubit and gate estimates. Identify recurring stable frontiers that may amortize state preparation. Establish

the crossover hypothesis before hardware availability.

12.6 Stage 5: governed institutional deployment

Only after usefulness is demonstrated should quantum-conditioned paths influence production analysis. Automatic authoritative write-back should remain prohibited in high-risk settings. Deployment should begin as a parallel advisory channel with human comparison against classical retrieval.

Institutional and Epistemological Implications

13.1 The unit of cognition becomes a distribution

Classical organizational procedure often asks for one recommendation. The Quantum Brain begins earlier, with a distribution over possible evidence routes. This changes the epistemic primitive:

single retrieval: find the answer path,

distributional cognition: construct a governed field of plausible and conflicting paths.

The final decision may still be singular. The process that supports it is plural.

13.2 A computational separation of inquiry and authority

Quantum sampling, LLM interpretation, expert validation, and institutional approval occupy different layers. This separation is not bureaucratic overhead. It prevents computational sophistication from masquerading as truth.

13.3 Memory as an endogenous topology

If the system writes approved candidate relations back into the graph, institutional memory becomes performative. The method used to search memory changes what memory later contains. Quantum-conditioned navigation makes this dynamic more visible because its probability distribution is an explicit object that can be logged and compared.

13.4 The value of disagreement

Contradiction is not noise to be eliminated. In finance, strategy, regulation, and science, unresolved contradiction may be a valuable state. A well-designed Super Brain preserves disagreement with provenance, identifies what evidence would resolve it, and prevents one persuasive narrative from becoming the graph's central hub merely through repetition.

Conclusion

The Quantum Brain makes sense, but not for the most common reason offered for it. A quantum computer does not allow an institution to inspect thousands of complete answers simultaneously. It allows carefully encoded amplitudes to evolve and interfere so that measurement produces a different distribution over computational states.

That capability becomes relevant to an ExoBrain when the computational states represent admissible reasoning paths, evidence nodes, or constrained context portfolios. The knowledge graph supplies structured external memory. The semantic compiler translates language into a computational specification. Quantum walks, amplitude amplification, QAOA, annealing, or quantum-inspired routines search the path space. LLM agents interpret the retained portfolio. Governance controls authority and memory.

The proposed Super Brain is therefore hybrid, plural, recursive, and governed. Its ambition is not merely to answer faster. It is to reduce premature cognitive collapse: to preserve strong support, strong contradiction, causal structure, precedent, weak signals, and nonlocal connections long enough for disciplined deliberation.

The validated notebook sequence shows both the promise and the discipline required. Quantum-conditioned routines changed the distribution of retrieved evidence, expanded context coverage, and changed the topology of the evolved graph. A small QAOA prototype closely approximated the exact portfolio optimum. But the classical baseline matched or exceeded the quantum-conditioned system on the declared end-to-end utility. No quantum advantage was demonstrated.

That conclusion is not a retreat. It is the foundation of a serious research programme. We now possess a falsifiable architecture, a claim ladder, exact baselines, a governance model, and an executable laboratory. The next question is no longer whether the phrase *Quantum Brain* is imaginable. It is whether a specific quantum-conditioned pipeline can cross the boundary from difference to usefulness and, eventually, from usefulness to computational advantage.

Final proposition

The deepest opportunity is not to build a quantum machine that “knows everything.” It is to build a governed cognitive system that can search a vast space of possible connections without losing provenance, plurality, humility, or human authority.

Compact Mathematical Architecture

A.1 State sequence

$$\begin{aligned}
 \Xi_t &= \text{Compile}(q_t, \mathcal{G}_t, \rho_t), \\
 \mathcal{G}_t^F &= \text{Frontier}(\mathcal{G}_t, \Xi_t), \\
 |\psi_0\rangle &= \text{Encode}(\mathcal{G}_t^F, \Xi_t), \\
 |\psi^\star\rangle &= U_{\kappa_t} |\psi_0\rangle, \\
 Z_t &= \text{Measure}_M(|\psi^\star\rangle), \\
 C_t^Q &= \text{Retain}_B(\text{Validate}(Z_t)), \\
 \mathcal{A}_t &= \text{Deliberate}(q_t, C_t^Q), \\
 \Delta\mathcal{G}_t^{\text{cand}} &= \text{Propose}(\mathcal{A}_t, C_t^Q), \\
 \mathcal{G}_{t+1} &= \mathcal{G}_t \oplus \text{Approve}(\Delta\mathcal{G}_t^{\text{cand}}).
 \end{aligned}$$

A.2 Minimum comparative estimand

For backend b , define end-to-end utility

$$U_b = w_a A_b + w_d D_b + w_c C_b + w_n N_b - w_e E_b - w_r R_b,$$

where A is answer validity, D is evidence diversity, C is contradiction coverage, N is nonlocal valid discovery, E is error or hallucination, and R is total resource cost. The usefulness estimand is

$$\Delta U = U_Q - U_C.$$

Quantum computational advantage requires not only $\Delta U > 0$, but a reproducible resource or scaling superiority under matched access assumptions.

Notebook-to-Theory Map

Notebook	Theoretical object	Principal outputs
NB17	Query compilation and explicit path space	Baseline graph, queries, 320-path catalogue, classical selectors, encoding manifest
NB18	Topology-sensitive amplitude evolution	Quantum and classical walk distributions, time sweep, measured portfolios, dephasing stress, navigation records
NB19	Constrained context portfolio	Exact feasible optimum, greedy and annealing baselines, QAOA statevector, approximation ratio, sample distribution
NB20	Evidence-bound deliberation	Six role contributions, recommendation, confidence, unresolved evidence, shock sensitivity, memoranda
NB21	Recursive governed evolution and claim discipline	Candidate registry, approval gate, evolved graphs, topology distance, noise benchmark, advantage claim ladder, audit ZIP

Minimum Audit Bundle

Each experimental run should retain:

1. graph version and cryptographic hash;
2. query and structured compilation record;
3. frontier nodes, edges, exclusions, and hash;
4. encoding specification and objective weights;
5. algorithm, backend, qubits, depth, shots, and noise model;
6. raw or summarized sample distribution;
7. classical baseline and resource accounting;
8. retained paths and evidence packets;
9. role contributions and synthesis;
10. candidate relations and governance decisions;
11. final graph diff;
12. claim status: difference, usefulness, or advantage.

Bibliography

- Andris Ambainis. Quantum walk algorithm for element distinctness. *SIAM Journal on Computing*, 37(1):210–239, 2007. doi: 10.1137/S0097539705447311.
- Panagiotis K. Barkoutsos, Giacomo Nannicini, Anton Robert, Ivano Tavernelli, and Stefan Woerner. Improving variational quantum optimization using CVaR. *Quantum*, 4:256, 2020. doi: 10.22331/q-2020-04-20-256.
- M. Cerezo, Andrew Arrasmith, Ryan Babbush, Simon C. Benjamin, Suguru Endo, Keisuke Fujii, Jarrod R. McClean, Kosuke Mitarai, Xiao Yuan, Lukasz Cincio, and Patrick J. Coles. Variational quantum algorithms. *Nature Reviews Physics*, 3:625–644, 2021. doi: 10.1038/s42254-021-00348-9.
- Andrew M. Childs and Jeffrey Goldstone. Spatial search by quantum walk. *Physical Review A*, 70(2):022314, 2004. doi: 10.1103/PhysRevA.70.022314.
- Andrew M. Childs, Richard Cleve, Enrico Deotto, Edward Farhi, Sam Gutmann, and Daniel A. Spielman. Exponential algorithmic speedup by a quantum walk. In *Proceedings of the Thirty-Fifth Annual ACM Symposium on Theory of Computing*, pages 59–68, 2003. doi: 10.1145/780542.780552.
- Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, and Jonathan Larson. From local to global: A graph RAG approach to query-focused summarization. 2024. URL <https://arxiv.org/abs/2404.16130>.
- Daniel J. Egger, Jakub Marecek, and Stefan Woerner. Warm-starting quantum optimization. *Quantum*, 5:479, 2021. doi: 10.22331/q-2021-06-17-479.
- Edward Farhi, Jeffrey Goldstone, and Sam Gutmann. A quantum approximate optimization algorithm, 2014. URL <https://arxiv.org/abs/1411.4028>.
- Austin Gilliam, Stefan Woerner, and Constantin Gonciulea. Grover adaptive search for constrained polynomial binary optimization. *Quantum*, 5:428, 2021. doi: 10.22331/q-2021-04-08-428.
- Lov K. Grover. A fast quantum mechanical algorithm for database search. In *Proceedings of the Twenty-Eighth Annual ACM Symposium on Theory of Computing*, pages 212–219, 1996. doi: 10.1145/237814.237866.
- Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yoon, Preslav Nakov, and Yu Su. HippoRAG: Neurobiologically inspired long-term memory for large language models. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- Xiaoxin He, Yijun Tian, Yifei Sun, Nitesh V. Chawla, Thomas Laurent, Yann LeCun, Xavier Bresson, and Bryan Hooi. G-retriever: Retrieval-augmented generation for textual graph understanding and question answering. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- Hsin-Yuan Huang, Michael Broughton, Masoud Mohseni, Ryan Babbush, Sergio Boixo, Hartmut Neven, and Jarrod R. McClean. Power of data in quantum machine learning. *Nature Communications*, 12:2631, 2021. doi: 10.1038/s41467-021-22539-9.

- IBM Quantum. How IBM will build a large-scale, fault-tolerant quantum computer, June 2025. URL <https://www.ibm.com/quantum/blog/large-scale-ftqc>. Accessed July 28, 2026.
- IBM Quantum. IBM to invest \$10 billion in quantum computing, June 2026. URL <https://www.ibm.com/quantum/blog/10-billion-investment-faq>. Accessed July 28, 2026.
- Andrew Lucas. Ising formulations of many NP problems. *Frontiers in Physics*, 2:5, 2014. doi: 10.3389/fphy.2014.00005.
- Jarrod R. McClean, Sergio Boixo, Vadim N. Smelyanskiy, Ryan Babbush, and Hartmut Neven. Barren plateaus in quantum neural network training landscapes. *Nature Communications*, 9:4812, 2018. doi: 10.1038/s41467-018-07090-4.
- Ashley Montanaro. Quantum algorithms: An overview. *npj Quantum Information*, 2:15023, 2016. doi: 10.1038/npjqi.2015.23.
- John Preskill. Quantum computing in the NISQ era and beyond. *Quantum*, 2:79, 2018. doi: 10.22331/q-2018-08-06-79.
- Jiashuo Sun, Chengjin Xu, Luming Tang, Saizhuo Wang, Chen Lin, Yeyun Gong, Heung-Yeung Shum, Jian Guo, and Daxin Zhu. Think-on-graph: Deep and responsible reasoning of large language model on knowledge graph. 2023. URL <https://arxiv.org/abs/2307.07697>.
- Ewin Tang. A quantum-inspired classical algorithm for recommendation systems. In *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing*, pages 217–228, 2019. doi: 10.1145/3313276.3316310.