

Beyond the Big Four

AI, Institutional Knowledge, and the Future Architecture
of Audit and Professional Services

From Generative AI to Institutional Intelligence—
and from Private Intermediation to Knowledge as a Public Good

Alejandro Reynoso

Honorary Fellow and External Lecturer
Cambridge Judge Business School
University of Cambridge

Research Note. This paper examines the potential structural impact of artificial intelligence on large audit and consulting organizations. Particular attention is given to institutional knowledge architectures, agentic systems, competition in professional services, and the possibility of treating portions of regulatory and assurance knowledge infrastructure as a public good.

The analysis is conceptual and is intended for research, educational, and policy discussion purposes.

August 2026

Executive Summary

The central argument. Artificial intelligence will do more than make auditors and consultants faster. When it is connected to persistent memory, validated institutional knowledge, provenance, analytical tools, specialized agents, and enforceable governance, it becomes part of the institution's operating architecture. The strategic asset is therefore not AI alone, but trusted institutional knowledge made computational.

The consequence for the Big Four. The same technology can strengthen and weaken the large professional networks. Their decades of accumulated judgments, methods, regulatory experience, and client history could become a formidable new moat if converted into governed computational capital. At the same time, AI may loosen the historic relationship between the scale of expertise and the scale of employment, allowing smaller firms to compete without reproducing the traditional professional pyramid.

The client-side challenge. Corporations can build second brains of their own. A persistent corporate knowledge system can remember why decisions were made, observe how they performed, and prevent each new advisory engagement from beginning with the expensive reconstruction of facts the client already possesses. External advisers will increasingly have to offer independence, comparison, investigation, or judgment that the client's own system cannot supply.

The public-policy opportunity. An independent Public Assurance Intelligence Utility could organize selected accounting, auditing, regulatory, enforcement, filing, and economic knowledge as shared digital infrastructure. By lowering the fixed informational cost of sophisticated assurance, such a utility could expand competition without attempting to manufacture another firm identical to the existing Big Four. It should support multiple independent assurance providers, not become a centralized government super-auditor.

The irreducible boundary. Machines may monitor transactions, retrieve precedent, test evidence, and expose disagreement. They cannot guarantee independence, skepticism, ethical courage, or accountability. Human professionals must continue to investigate anomalies, adjudicate uncertainty, certify conclusions, and bear responsibility for them. The future of assurance is therefore computational in its reach but human in its ultimate authority.




Figure 1: Visual overview of the paper’s argument, institutional architecture, and possible futures for professional services.

1. Introduction

For more than a century, the modern professional firm has sold something that is difficult to see. Its offices, reports, and spreadsheets are visible enough, as are the armies of graduates who enter each year and the partners whose signatures ultimately carry legal and reputational weight. The real product, however, has always been organized confidence. A board hires an auditor or consultant because the firm appears able to find the relevant facts, summon the right expertise, compare the present problem with a long history of similar problems, and produce a conclusion that other powerful institutions will accept.

Generative artificial intelligence entered this world looking, at first, like an unusually capable junior colleague. It could read a contract in seconds, summarize a filing, draft a memorandum, search an accounting standard, or construct the first version of a model. These are meaningful gains. They reduce drudgery, shorten turnaround times, and allow individual professionals to cover more ground. In industries whose economics have long been tied to the hours consumed by information processing, even that limited transformation is destabilizing.

But productivity is only the opening chapter. The deeper change begins when the system does not forget at the end of the conversation; when it can distinguish a binding standard from an internal opinion; when it can trace a conclusion to the evidence, calculation, and precedent that produced it; when it can call reliable analytical tools rather than improvise an answer; and when its actions are bounded by permissions and reviewed by people whose responsibility cannot be delegated to a machine. At that point, AI is no longer simply helping a professional perform a familiar task. It is changing the way the organization remembers, coordinates, and decides.

This paper examines that second transformation. Its subject is not whether software can make a consultant quicker or an audit cheaper. It is what happens when the knowledge on which professional intermediation depends becomes persistent, structured, machine-readable, and partly executable. The question is larger than technology. It reaches into the design of firms, the economics of expertise, the concentration of audit markets, the responsibilities of regulators, and the institutional foundations of trust.

2. The Economic Logic of the Traditional Professional Firm

The traditional professional partnership arose as an answer to a brutal coordination problem. Expertise took years to acquire and was unevenly distributed. Information moved slowly. The memory of an engagement lived largely in the minds of the people who had worked on it, in paper files, and later in digital repositories that were comprehensive but rarely easy to navigate. Complex work had to be broken into pieces, assigned to teams, reviewed through layers, and finally assembled into a judgment by a relatively small group of senior professionals.

The familiar pyramid was therefore not merely a cultural habit. It was an information-processing machine built out of people. Junior staff collected evidence and performed repetitive tests. More

experienced professionals identified exceptions and interpreted rules. Managers coordinated the work, and partners reconciled the analysis with commercial reality, professional standards, and the firm's appetite for risk. Each layer compressed and refined the information moving upward. The system could be slow and costly, but it allowed an institution to turn thousands of fragments into a defensible conclusion.

Scale mattered because knowledge was attached to human beings. A firm that wanted to cover more industries, jurisdictions, tax regimes, technologies, and financial instruments needed more people who had encountered those things. It also needed a larger training system, deeper review structures, and a wider network of specialists who could be called when an engagement crossed an unfamiliar line. Organizational scale and knowledge scale grew together, reinforcing the dominance of global firms able to carry the fixed cost of that apparatus.

Artificial intelligence does not abolish the need for expertise, but it may alter the machinery through which expertise is gathered and distributed. If the lessons once trapped in personal memory and scattered files can be made available across an institution—with their context and provenance intact—then the capacity to know more need no longer rise in a straight line with the number of people on the payroll. That is the fault line running beneath the contemporary debate.

3. The Big Four Are Already Moving

None of the major audit networks is standing still. Deloitte, PwC, EY, and KPMG have invested heavily in cloud infrastructure, audit analytics, generative AI, and the emerging class of systems commonly described as agentic. Deloitte has added generative and agentic capabilities to its Omnia audit environment. PwC has connected generative systems to its professional research and has discussed the development of AI-enabled audit platforms. EY and KPMG have likewise placed AI inside broader programs of assurance technology and workflow modernization.

Regulators have begun to follow the technology into the audit room. The UK Financial Reporting Council has addressed the use of generative and agentic AI in audit; the International Auditing and Assurance Standards Board has made technology part of its strategic agenda; and the US Public Company Accounting Oversight Board has examined how generative AI is entering auditing and financial reporting. These developments matter because audit technology does not operate in an ordinary market. A system may be ingenious and still be unusable if it cannot satisfy requirements of evidence, documentation, independence, inspection, and accountability.

For the moment, much of the industry's activity remains organized around a recognizable model: the human professional is still the center of the workflow, while AI searches, drafts, summarizes, and suggests around the edges. That model is already changing the work. Yet it leaves the deeper structure of the firm mostly untouched. The next frontier begins when the organization treats AI not as a collection of personal assistants, but as an institutional layer through which its knowledge can be preserved, tested, and deployed.

4. Generative AI Is Not Institutional Intelligence

A foundation model can appear knowledgeable in almost any room. It can explain an accounting rule, compare regulatory regimes, review a contract, or sketch the logic of a valuation. That breadth is impressive, but a professional institution depends on a different kind of knowing. It must remember what happened in this company, on this engagement, under this version of the standard, with this evidence, and before this committee. It must know not only the answer that was accepted, but which alternatives were considered, why they were rejected, and who assumed responsibility for the decision.

A general model does not possess that history merely because it can speak fluently about the subject. It does not automatically know that a control failed two years ago, that a technical panel rejected a superficially similar interpretation, or that the decisive piece of evidence arrived after the initial memorandum was drafted. Nor can it be trusted to separate law, firm methodology, client assertion, historical practice, and its own inference unless the system around it makes those distinctions explicit.

Institutional intelligence therefore requires more than reasoning. It requires memory that persists, knowledge that has been curated, context that travels with the record, provenance that allows every important claim to be traced, and governance that determines what the system may do and who must answer for it. The foundation model may be the most conspicuous component, but it is only one component. The institution is created by the architecture around it.

5. The Institutional Second Brain

The phrase *institutional second brain* risks sounding more mystical than the idea deserves. In practice, it describes a set of disciplined connections. The institution must connect general reasoning to its own validated knowledge; link documents to the people, rules, evidence, and events to which they refer; preserve the origin and status of every important claim; route numerical work to reliable tools; and place consequential actions inside a visible chain of human authority. What makes the system intelligent is not any single database or model, but the way these parts work together.

5.1 Foundation Intelligence

At the outer edge sit foundation models capable of language, reasoning, code, mathematics, and the interpretation of images and documents. Professional firms will naturally compete to use the best available systems, but they should resist building their entire strategy around a particular model. Models will improve, prices will change, vendors will rise and fall, and different tasks will favor different technical approaches. A durable institution should be able to replace a model without losing its memory. Its strategic asset is the governed knowledge environment into which models are invited, not the temporary identity of the model at the door.

5.2 Curated Institutional Knowledge

Behind that door lies the institution’s intellectual capital. Some of it is public and authoritative: accounting and auditing standards, legislation, regulatory guidance, and enforcement decisions. Some is proprietary: methodologies, technical interpretations, industry research, templates, analytical models, and the accumulated record of previous engagements. Still more exists in the neglected border country between formal rule and lived experience—the exceptions that forced a change in approach, the judgments that survived review, and the warnings that later proved prescient.

Simply collecting these materials is not enough. Professional firms already possess enormous repositories. The difficult work is to decide what is valid, current, superseded, disputed, confidential, or relevant only under particular circumstances. Curation is not clerical housekeeping; it is the process by which an archive becomes usable knowledge.

5.3 Knowledge Graph

The decisive move is from storage to relationship. An audit issue rarely lives inside one document. A transaction points to an accounting treatment; the treatment depends on a standard and an interpretation; the interpretation depends on facts and evidence; the evidence is produced by systems and controls; the controls shape the risk assessment; and the risk assessment determines the audit procedure and, ultimately, the conclusion. A conventional repository may contain every document in that chain while leaving the chain itself invisible.

A knowledge graph makes those relationships explicit. It allows the institution to ask not merely for every file containing a phrase, but for every conclusion that relied on a given control, every engagement affected by a revised interpretation, or every valuation whose assumptions were drawn from the same market source. This is how memory becomes analytical rather than merely archival.

5.4 Provenance

In professional work, an answer without a lineage is often no answer at all. The system must be able to show whether a proposition comes from binding regulation, an accounting or auditing standard, the firm’s own methodology, evidence supplied by a client, an external study, a historical precedent, a model’s inference, or a human judgment made under uncertainty. These categories cannot be blended into one smooth paragraph and treated as equivalent.

Provenance is therefore the line between assistance and authority. It enables a reviewer to reconstruct the path from source to conclusion and to see where interpretation entered. Without it, AI’s fluency becomes a professional hazard: the system can make weak evidence sound settled and old assumptions sound current. With it, the system becomes open to challenge, correction, and inspection.

5.5 Deterministic Tools

The architecture must also know when not to use a language model. Reconciliations, statistical tests, sampling, transaction matching, valuation calculations, and many accounting computations should be performed by deterministic tools whose methods and outputs can be reproduced. The model's role is to understand the question, select the appropriate tool, interpret the result, and explain its limits. Allowing a probabilistic language system to improvise arithmetic that a tested engine can perform exactly is not innovation. It is a failure of design.

5.6 Specialized Agents

Specialized agents can then work across this combination of knowledge and tools. One may examine the accounting treatment; a second may test the consistency of the evidence; another may search for fraud indicators, evaluate controls, review a valuation, or challenge a going-concern assumption. Their value lies in sustained, bounded attention. Each can pursue a defined question, preserve its work, and pass the result to another part of the system.

The word *agent*, however, should not become an alibi for unrestricted autonomy. In an assurance environment, an agent must have a defined mandate, limited permissions, known sources, and a clear escalation path. It should be able to stop when evidence is missing, expose disagreement rather than conceal it, and return consequential decisions to accountable professionals.

5.7 Governance

Governance binds the architecture together. Every significant process should leave a record of who initiated it, what information was accessed, which model and tools were used, what assumptions were introduced, what evidence was considered, what conclusion emerged, and which people reviewed and approved it. The audit trail must include the machine as well as the humans.

These requirements may seem to slow the technology, but in professional services they are what make the technology usable. Trust does not arise because a system is confident or convenient. It arises because the system can be questioned and because a person or institution remains answerable when the result is wrong. Once those conditions are present, the second brain is no longer a chatbot attached to a document store. It has become part of the operating constitution of the firm.

6. A Conceptual Architecture

Figure 2 presents the architecture as a sequence of layers, although in practice information would move repeatedly in both directions. A model may begin with a question, retrieve a relevant standard and engagement history, follow their relationships through the knowledge graph, invoke an analytical tool, and assign a specialized agent to test the result. The governance layer surrounds the entire process: it does not merely appear at the end.




Figure 2: Conceptual architecture of an institutional second brain.

The diagram is intentionally model-agnostic. Its central assumption is that reasoning systems will change faster than the institutions that use them. Knowledge, provenance, permissions, and accountability must therefore survive technical substitution. A firm that preserves those elements can benefit from better models as they appear. A firm that confuses a vendor relationship with an institutional knowledge strategy may discover that it has automated fragments of work without building a durable capability.

7. Why This Challenges the Traditional Consulting Model

Consulting has long converted labor into revenue with unusual directness. A problem is divided among people, their time is recorded, and the client is charged for the hours required to move from raw information to recommendation. The leverage model rewards the firm for placing a broad base of junior labor beneath a narrow group of senior judgment. Even fixed-fee projects are usually priced against an internal estimate of how many people, at which levels, will be needed for how long.

AI places pressure on the middle of that equation. Research, comparison, first-pass analysis, document review, and drafting can often be completed with far fewer human hours. The result does not eliminate the need for judgment; it changes the route by which the team reaches the point where judgment is needed. Data can move through institutional knowledge, analytical tools, and specialized agents before arriving at a senior professional for challenge and decision. Work that once

justified several layers of coordination may be performed by a smaller group with a much larger computational reach.

This creates a commercial dilemma. If a firm charges by the hour, efficiency reduces the apparent volume of what it has to sell. If it conceals the efficiency and preserves the old price, clients will eventually notice that the underlying cost has changed. The durable response is to move the basis of value away from the consumption of labor and toward knowledge, trust, assurance, judgment, and outcomes. The firm of the future will not prosper merely because it gives a client access to well-informed people. It will prosper if it can deliver institutional intelligence that is demonstrably reliable and accept responsibility for the result.

8. Consulting and Audit Face Different Disruption Paths

Consulting and audit will not travel this road at the same speed. A company can postpone a strategy project, build an internal team, or decide that an outside recommendation is not worth buying. It usually cannot opt out of a statutory audit. That legal requirement creates a formidable moat around the audit profession, and the signature of an independent auditor retains a social function that no productivity tool can simply erase.

The moat protects the mandate, however, not every task performed inside it. An audit combines evidence collection, testing, interpretation, judgment, and certification. Machines can already assist with the first three, and in well-structured environments they may perform a growing share of them. They can also inform judgment by finding patterns and precedents that no individual could hold in mind. What remains irreducible is the institutional act of standing behind a conclusion: challenging management, weighing ambiguity, deciding whether the evidence is sufficient, and accepting the legal and professional consequences of the opinion.

Audit is therefore less likely to vanish than to change shape. The workforce may become smaller and more technically specialized. The routine processing that once trained generations of junior auditors may decline, forcing firms to invent new routes by which novices acquire judgment. The center of professional value will move toward skepticism, investigation, systems governance, and certification—the parts of the audit that matter most when the answer is contested.

9. The Second-Brain Moat

Buried inside the Big Four is an asset more valuable than any single model: the residue of millions of professional encounters. Their firms have watched accounting treatments succeed and fail, controls hold and collapse, valuations survive scrutiny and unravel under it. They have seen tax structures tested, cyber incidents spread, transactions change shape under regulatory pressure, and governance failures reveal themselves first as small inconsistencies. No new entrant can quickly recreate that history.

Yet possession is not the same as use. Much of this knowledge is fragmented across engagement files,

methodology portals, specialist databases, archived memoranda, and the memories of people who may have moved to another office or retired. Confidentiality and legal restrictions properly limit how engagement information can be reused. Even where access is permitted, the decisive lesson may be buried in a document whose title gives no clue to its importance.

The strategic opportunity is to transform that tacit and scattered experience into governed computational capital. This does not mean pouring every old file indiscriminately into a model. It means extracting reusable patterns, preserving necessary boundaries, recording the authority and context of each lesson, and making the result available to the right people and systems at the right time. If the firms can accomplish that, their age becomes an advantage: decades of experience can be made present at the moment of decision. The institutional second brain may then become their strongest defense against commoditization.

10. The Client-Side Paradox

The same architecture that can deepen the moat around a professional firm can also help a client climb over it. A multinational company possesses its own dense history of accounting decisions, tax positions, contracts, controls, treasury policies, operational exceptions, and regulatory encounters. Until recently, that history was often no easier for the company to navigate than it was for an outside adviser. Each new project began with familiar rituals: interviews, document requests, data rooms, and the costly reconstruction of facts that the organization technically already knew.

A corporate second brain changes that balance. Unlike the consulting team, it does not leave when the engagement ends. It can preserve the reasoning behind a decision, observe how the decision performs, and connect later evidence to the original assumptions. Over time, the system can know more about the corporation's recurring problems than an external adviser could learn during a conventional engagement.

The consequence is not the end of outside advice. External experts still bring independence, comparative experience, scarce specialties, and the capacity to challenge an organization's settled beliefs. But clients may become less willing to pay outsiders to rediscover their own history. The consultant will have to contribute something the corporate intelligence system does not already contain: a genuinely new comparison, an independent challenge, a difficult judgment, or a credible commitment to stand behind the answer.

11. The Regulatory Second Brain

Public authorities face the same problem of fragmented knowledge, but at a larger scale and with higher stakes. The intelligence of a regulator is scattered across statutes, rules, professional standards, supervisory guidance, enforcement actions, corporate filings, inspection findings, market data, and the institutional memory of civil servants. Each source is valuable; their connections are where warning signs often emerge.

A regulatory second brain could connect those sources while preserving their legal status and restrictions. It could identify companies whose disclosures drift from their peers, trace repeated control weaknesses across reporting periods, compare an auditor's current judgments with earlier inspection findings, or show how a new market shock affects clusters of firms sharing similar assumptions. Such a system would not need to wait for the annual cycle to assemble a picture from scratch. It could update its view as new filings, transactions, and supervisory evidence arrive.

That movement from periodic supervision toward continuous computational supervision would be significant. It could allow regulators to direct scarce human attention toward the cases where signals converge, rather than spreading it evenly across an enormous population. It would also create new risks: an error in a shared model could affect many decisions, surveillance could outrun legal authority, and opaque risk scores could harden suspicion into an administrative fact. The public system would therefore require even stronger provenance, transparency, and routes of appeal than its private counterpart.

12. The Public Assurance Intelligence Utility

A more ambitious possibility is an independent Public Assurance Intelligence Utility, or PAIU. The analogy is not to a ministry that issues the final audit opinion, but to infrastructure on which many participants can build. The utility would organize selected public knowledge, maintain validated analytical resources, and offer standard interfaces to auditors, companies, regulators, researchers, and investors. It would reduce the waste involved when every institution separately collects and structures much of the same public information.

The PAIU would not replace statutory auditors. Its purpose would be to improve the common informational ground on which auditors compete and regulators supervise. The final judgment would remain with an independent provider, and the final enforcement power with the relevant public authority. The utility would make the underlying knowledge infrastructure more accessible, transparent, and capable.

12.1 Public Knowledge Layer

Its public knowledge layer could connect accounting and auditing standards with company law, securities regulation, regulatory interpretations, enforcement precedents, historical filings, economic data, and industry classifications. Properly anonymized inspection findings could reveal recurring weaknesses without exposing confidential engagement details. Versioning would allow a user to reconstruct what the authoritative framework said at the time a decision was made, rather than silently applying today's rule to yesterday's facts.

12.2 Analytical Layer

Above that foundation, a validated analytical layer could provide accounting engines, statistical tests, benchmarking, network analysis, anomaly detection, fraud indicators, and going-concern analytics. These tools would not declare a company sound or an audit sufficient. They would produce reproducible signals that competing assurance providers could interpret, challenge, and supplement with their own methods.

12.3 Governance Layer

The governance layer would determine whether the utility deserved public trust. Its sources and methodologies would need documented provenance and version control. Models would require independent evaluation, security testing, and continuing review. Material changes could not be smuggled into production as technical updates; they would need a visible process and, where appropriate, public consultation. Institutions affected by an output would need a way to challenge the data or method that produced it. Academic and civil- society scrutiny would be a feature of the design, not an intrusion.

12.4 Interface Layer

Finally, standardized interfaces would allow different participants to use the infrastructure without forcing them into a single system or interpretation. A small audit firm might use the public knowledge graph while applying its own sector model. A university might test whether a risk indicator performs as claimed. A corporation might compare its disclosures with market patterns, while a regulator receives a more detailed view under lawful authority. The governing principle would be common knowledge infrastructure combined with competing assurance providers, not a single machine issuing a single official truth.

13. Knowledge as a Public Good

Much of the intellectual foundation of audit is public already. The standards can be read, regulations consulted, filings downloaded, enforcement decisions studied, economic series analyzed, and decades of academic work reviewed. What remains costly is the labor of collecting these materials, keeping them current, reconciling their terminology, mapping their relationships, and turning them into a working system. Large firms can spread that fixed cost across vast networks. Smaller firms cannot.

AI changes the economics of the organizing layer. Once public knowledge has been structured and supplied with reliable provenance, many users can draw on it without rebuilding the same machinery. The benefit is not merely administrative efficiency. Shared infrastructure could lower the minimum scale at which an assurance provider can compete with sophistication. A challenger would still need excellent people, sector expertise, secure systems, insurance, licenses, and the confidence of

boards and regulators. It would not, however, have to reproduce every element of the informational backbone developed by firms over decades.

This suggests a different approach to audit-market competition. Policy has often searched for a fifth firm large enough to resemble the existing four. That ambition collides with the very economics that produced concentration in the first place: global reach, specialist depth, technology costs, regulatory demands, and the liability associated with large public companies. Public knowledge infrastructure would attack the problem from below by reducing one of the fixed costs of credible entry. Rather than manufacture another pyramid, it could make the pyramid less necessary.

14. Why the State Should Not Become the Ultimate Auditor

The case for public infrastructure is not a case for a government super-auditor. A single state system that gathers the data, chooses the model, interprets the evidence, certifies the accounts, and enforces the result would concentrate too much epistemic and coercive power in one place. Political pressure could shape supposedly technical judgments. A cyberattack or model defect could become systemic. Legitimate differences of interpretation might be treated as deviations from the machine's official answer.

Institutional design should therefore separate the provision of knowledge infrastructure from analysis, judgment, certification, and enforcement. Public bodies may maintain authoritative sources and common tools without deciding every case. Private and public assurance providers may analyze the same evidence through different methods. Independent auditors may certify conclusions, while courts and regulators retain distinct mechanisms for enforcement and appeal. The frictions among these institutions are not waste to be engineered away. They are safeguards against the confidence of a single system.

15. A Federated Assurance Architecture

A better arrangement would be federated. Common public knowledge could feed multiple independent systems operated by audit firms, regulators, exchanges, universities, and other qualified bodies. Those systems would not be expected to agree automatically. One might regard a transaction as adequately supported, another might flag it for investigation, and a third might reject the accounting treatment. Their disagreement would be information.

Today, differences of professional judgment are often discovered late, after each institution has conducted much of its analysis in isolation. A federated architecture could reveal the point of divergence earlier: perhaps the systems used different market data, weighted a precedent differently, or disagreed about the reliability of a control. Human professionals could then adjudicate the actual conflict rather than argue across opaque conclusions.

Such pluralism would make the system less tidy and more resilient. The objective of assurance is not to eliminate all disagreement but to ensure that consequential disagreement becomes visible,

reasoned, and answerable. A market governed by several inspectable systems is safer than one governed by an oracle, whether that oracle belongs to a technology company, a global audit network, or the state.

16. The Future Auditor as Governor of Machines

The auditor formed by this system will look different from the professional formed by the old pyramid. Much less time may be spent copying evidence, reconciling spreadsheets, selecting routine samples, searching standards, or drafting memoranda whose structure barely changes from year to year. Machines are well suited to the patience and repetition those tasks demand.

The human role will move toward the moments when the process becomes uncertain or consequential. Auditors will need to understand why a model produced a signal, whether the data deserve trust, which causal story best explains an anomaly, and when management's account has become too convenient. They will govern permissions, test the behavior of agents, challenge embedded assumptions, investigate exceptions, and decide what must be escalated. Above all, they will remain responsible for judgments that affect companies, investors, employees, and markets.

Calling the future auditor a governor of machines is not rhetorical promotion. Governance is a demanding technical and ethical craft. It requires enough understanding of the system to know where it can fail, enough professional experience to recognize what the system has missed, and enough independence to resist both management and the seductive confidence of automation.

17. From Periodic Audit to Continuous Assurance

Traditional audit is periodic because human attention is scarce. Teams arrive, gather evidence, select samples, perform tests, and compress a year of corporate activity into an opinion issued after the period has ended. The business itself never pauses in this way. Transactions, journal entries, cash movements, control events, valuations, cyber incidents, and related-party dealings continue to accumulate every day.

As more evidence becomes machine-readable, monitoring can move closer to the rhythm of the activity being monitored. Agents could examine complete populations of transactions, watch for changes in control behavior, compare disclosures across systems, and alert auditors when several weak signals begin to form a pattern. Human investigation could then be concentrated on the exceptions carrying the greatest risk, rather than on a periodic sample selected largely because it is possible to inspect by hand.

Continuous assurance should not be confused with constant certainty. Data can be incomplete, controls can be gamed, and an automated test can reproduce the same mistaken assumption millions of times. A poorly governed system does not merely make errors faster; it can make them uniform and harder to notice. Continuous monitoring will improve assurance only if the rules, models, data, and escalation paths are themselves continuously challenged.

18. The New Challenger

The most dangerous challenger to a Big Four firm may not resemble a fifth member of the club. Imagine instead a regulated organization with a few hundred exceptional accountants, engineers, quantitative specialists, lawyers, and cybersecurity experts. Its institutional memory resides in a carefully governed knowledge graph. Specialized agents conduct much of the initial analysis. Every step has provenance, methodologies are tested continuously, and human experts focus on the cases where judgment is genuinely required.

Such a firm would still face imposing barriers. It would need authorization, insurance, capital, geographic reach, secure access to client systems, and a reputation strong enough for audit committees to entrust it with consequential work. It would also need a credible answer to the liability created by failure. Technology does not make those institutions appear overnight.

What technology changes is the relationship between expertise and employment. The challenger might achieve the analytical reach of a much larger organization without reproducing its headcount or hierarchy. Its memory would survive turnover, its methods could be reproduced across engagements, and new knowledge could propagate without waiting for thousands of people to attend a training course. The competitive unit would no longer be the professional alone, but the professional amplified by an institutional architecture.

19. The Defensive Advantages of the Big Four

The incumbent firms are hardly defenseless. Their names carry weight in boardrooms and capital markets. They hold licenses, maintain regulatory relationships, operate global networks, absorb enormous liability, train large populations of professionals, and possess specialists across industries and jurisdictions. Trust in audit is partly technological, but it is also historical and institutional. A new entrant cannot download it.

The central question is whether the Big Four use these advantages to build a new architecture or merely to protect the old one. If AI is deployed chiefly to produce the same report with fewer hours, the firms will gain efficiency and perhaps a temporary expansion in margin. Clients and competitors will eventually claim much of that gain. If, instead, the firms turn decades of professional experience into a persistent and governed intelligence system, they may create an asset that grows more useful with every engagement.

The distinction is between doing the old work faster and changing what the firm is. Productivity matters; transformation changes the basis of competition. The incumbents possess the raw material for the second path, but their size can make the conversion difficult. Knowledge is divided by national network, service line, client confidentiality, technology platform, and professional hierarchy. The very structure that accumulated the experience can obstruct its use.

20. From AI Strategy to Knowledge Strategy

For that reason, the question placed before a board should not begin and end with, “What is our AI strategy?” The phrase invites a list of vendors, pilots, and use cases. It encourages the institution to think about the technology it can buy rather than the knowledge it must govern.

The more important question is what the institution knows and what happens to that knowledge over time. Where does it reside? Which part is authoritative, which is provisional, and which has quietly become obsolete? Who may use it, and under what obligations? Can the firm trace a conclusion back to its evidence? Can a later event revise the record without erasing the reasoning that came before? When two experts disagree, does the system preserve the disagreement or force an artificial consensus? What disappears every time an experienced employee walks out of the door?

Those are questions of knowledge strategy. They determine how models and agents should be used, what data architecture is required, and where human approval belongs. A firm that answers them can change models without losing itself. A firm that ignores them may acquire an impressive collection of AI tools while remaining institutionally forgetful.

21. Knowledge as the Scarce Asset

As powerful models become widely available, generic intelligence is likely to become less distinctive. Competing firms will be able to rent similar reasoning capacity from the same vendors or switch among several capable systems. The model may remain essential while ceasing to be a moat, much as cloud computing is essential to modern business without usually constituting a unique strategy.

Scarcity will migrate toward trusted knowledge and the institutional capacity to use it. A current, well-structured record of how a rule has been interpreted across industries is more valuable than a model that can merely recite the rule. A provenance system showing why a conclusion was reached is more defensible than a fluent answer without sources. A governance process that knows when to stop a machine and summon a human is more consequential than another increment of benchmark performance.

The emerging moat is therefore a compound asset: trusted knowledge, computational intelligence, governance, and human judgment operating together. None is sufficient alone. Knowledge without reasoning is an archive; reasoning without provenance is a risk; governance without capable tools is bureaucracy; and human judgment without access to the relevant evidence is intuition under pressure. The institutional second brain is strategic precisely because it joins these elements.

22. Three Possible Futures

The technology does not dictate a single destination. At least three broad futures are plausible, and elements of all three may coexist across countries and service lines.

22.1 Scenario I: AI-Enhanced Big Four

In the first future, the Big Four automate existing activities while preserving the broad architecture of the partnership. Research, testing, drafting, and review become faster; the pyramid grows thinner; and firms initially capture part of the efficiency in higher margins. Over time, clients demand lower prices and competitors adopt similar tools. AI becomes pervasive but not revolutionary—an important production technology operating inside a familiar business model.

22.2 Scenario II: Cognitive Big Four

In the second, the firms become cognitive institutions. Their accumulated experience is converted into computational capital; knowledge graphs become strategic assets; and specialized agents operate throughout professional workflows. Smaller groups of highly skilled people govern systems with far greater analytical reach. Revenue shifts away from the simple sale of hours and toward access to knowledge infrastructure, continuing assurance, difficult judgment, and outcomes for which the firm is willing to be held accountable. This is the strongest long-run strategy available to the incumbents, but it requires them to reorganize around knowledge rather than merely attach technology to existing service lines.

22.3 Scenario III: Assurance as Infrastructure

In the third future, assurance itself begins to resemble infrastructure. Regulators, exchanges, professional bodies, universities, and market participants build shared knowledge systems. Continuous computational supervision expands, and private auditors interact with a common informational foundation. Competition moves toward the qualities that cannot simply be copied from the shared layer: sector knowledge, investigative ability, specialized models, professional independence, and the credibility of certification.

This arrangement would not abolish private audit. It would change the division of labor between public knowledge and private judgment. The informational ground would become more common, while interpretation and responsibility would remain plural. The public would not be asked to trust a single state answer; it would gain a clearer view of the evidence and methods behind competing professional answers.

23. The Industrial Organization of Knowledge

The larger economic story is about the industrial organization of knowledge. The twentieth-century professional firm coordinated expensive expertise, fragmented information, human memory, and labor-intensive analysis through hierarchy. Its size was not an accident; size was the technology available for bringing many forms of knowledge to bear on one problem.

AI alters each term. Search makes more of the archive reachable. Persistent memory allows

experience to accumulate beyond the tenure of an individual employee. Computational analysis reduces the labor required to test and compare. Agents can coordinate bounded tasks without reproducing every managerial layer. Some forms of expertise that once existed only as scarce human practice can be partially encoded, evaluated, and distributed.

The consequence is not that employment becomes unnecessary. It is that knowledge scale may no longer require employment scale in the same proportion. A smaller institution may command a much wider field of expertise, while a large institution that fails to organize its knowledge may receive surprisingly little advantage from its headcount. That shift could matter more to the structure of professional services than the spectacle of generative AI itself.

24. The Irreducible Role of Human Judgment

There is an essential boundary to this argument. A second brain can remember more than any partner. A powerful model can reason across more documents than any engagement team. Neither capacity guarantees wisdom, independence, skepticism, or courage. A system trained perfectly on an institution's history may reproduce the institution's blind spots with extraordinary efficiency.

Professional judgment is not a decorative human layer placed on top of a completed machine analysis. It is the capacity to decide what the formal process has failed to see, to recognize when the available evidence is shaped by power, and to act when the commercially easy answer is professionally indefensible. It includes the willingness to tell a client, a senior partner, or a regulator that the prevailing interpretation is wrong.

Nor is judgment purely individual. Institutions create or destroy the conditions under which people can exercise it. Incentives, promotion systems, consultation procedures, independence rules, and legal accountability shape whether skepticism survives contact with commercial pressure. The final architecture of trust must therefore remain human and institutional, even when much of the underlying analysis becomes computational.

25. Conclusion

The Big Four are unlikely to disappear in a sudden flash of generative intelligence. They are protected by regulation, reputation, relationships, accumulated expertise, and the stubborn fact that markets need institutions willing and able to accept responsibility. But survival is not the same as continuity. The firms can remain powerful while becoming substantially different organizations.

Generative AI reduces the labor required to produce analysis. Agentic systems reduce the coordination required to move work through a hierarchy. Institutional second brains reduce the loss created by fragmented memory. Knowledge graphs make connections visible across silos, and continuous monitoring challenges an assurance model built around periodic human sampling. Each technology removes a small piece of the economic logic that once made the giant professional pyramid seem inevitable.

The incumbents therefore face a double possibility. AI may become the greatest productivity technology they have ever adopted, allowing them to deliver better work with fewer routine hours. The same technology may also commoditize part of the intellectual intermediation on which their business was built. If clients can preserve their own institutional memory, challengers can draw on shared knowledge infrastructure, and regulators can supervise more continuously, the mere possession of a large workforce will carry less strategic weight.

The decisive asset of the coming era is not artificial intelligence in the abstract. It is institutional knowledge made computational: knowledge linked to evidence, capable of being reasoned over, clear about its provenance, constrained by governance, and completed by human judgment. The winning organization may not own the largest language model or employ the most people. It may be the institution that can connect those elements with the greatest reliability and earn the right to be trusted when they point in different directions.

The same principle applies beyond the firm. Corporate, audit, and regulatory second brains may increasingly meet across governed interfaces, each carrying a different history and responsibility. Machines may monitor the routine, surface anomalies, reconstruct precedent, and expose disagreement. People will still have to investigate what matters, decide what the evidence means, certify the conclusion, and bear the consequences.

That is a more profound future than one in which an auditor writes a memorandum faster. It is the possible emergence of a new institutional architecture of trust—one built not on the fantasy that machines will remove uncertainty, but on the harder achievement of making knowledge, reasoning, disagreement, and responsibility visible to one another.

Selected References

- [1] Deloitte. *Deloitte Expands AI Capabilities in Omnia Global Audit Platform*. Deloitte, 2025. Available at: <https://www.deloitte.com/us/en/about/press-room/deloitte-expands-AI-capabilities-in-omnia-global-audit-platform.html>
- [2] Financial Reporting Council. *Artificial Intelligence in Audit: Guidance on Generative and Agentic AI*. FRC, London, 2026. Available at: <https://www.frc.org.uk/library/standards-codes-policy/audit-assurance-and-ethics/guidance/ai-in-audit/>
- [3] Financial Reporting Council. *Developments in Audit Market Competition and Public Interest Entity Audit*. FRC, London.
- [4] International Auditing and Assurance Standards Board. *Technology and the Future of Audit and Assurance*. IAASB. Available at: <https://www.iaasb.org/focus-areas/technology>
- [5] Public Company Accounting Oversight Board. *Staff Observations from Outreach on the Use of Generative Artificial Intelligence in Audits and Financial Reporting*. PCAOB, Washington, D.C.
- [6] PwC. *Audit Quality Report 2025*. PricewaterhouseCoopers, 2025.

- [7] EY. *Audit Quality Report*. Ernst & Young.
- [8] UK Government. *Restoring Trust in Audit and Corporate Governance*. Department for Business and Trade.
- [9] House of Commons Library. *Auditing Reforms and Competition in the UK Audit Market*. UK Parliament.
- [10] Financial Reporting Council. *Innovation and Improvement Hub*. FRC, London.
- [11] “The (Overdue) Collapse of Bullsh*t Companies.” YouTube video used as the initial motivation for this research note. Available at: <https://youtu.be/ohCAPyqlJpo>