

UNDERSTAND | BUY | BUDGET | GOVERN

How to BuyAI

A CEO's Guide to What You Are Buying, How to Budget It,
and How to Decide What Is Worth Funding

For CEOs, Entrepreneurs, Financial Practitioners & Boards of Directors

For decades, executives have bought technology with a familiar mental model: software licenses, servers, cloud capacity, implementation projects, and IT budgets. AI disrupts that comfort. A vendor may be selling a model, access to a model, a copilot, an agent, a workflow, a vector database, a private knowledge system, or an entire operating architecture—and all of these may be described simply as “AI.”

This guide starts from the non-specialist executive's real question: **What exactly am I buying, how will I be charged for it, what else must I budget around it, and what evidence should I require before approving the investment?**

Alejandro Reynoso

August 2026

Executive Practitioner Guide / Applied AI Architecture

Corporate Investment & Governance

A practical framework for understanding AI products, pricing, total cost of ownership, buy-versus-build choices, investment approval, operating economics, and board oversight.

Reader’s Map

The executive question

“I am not an AI specialist. My company has bought technology for years. What is different now—and how do I make a rational investment decision?”

The purpose of this guide is not to turn the CEO into a machine-learning engineer. It is to provide enough architectural and economic understanding to ask the right questions of management, vendors, advisers, and the board.

Section	The question it answers
1. Why AI feels different	Why traditional technology procurement instincts are necessary but no longer sufficient.
2. What are you buying?	The layers hidden behind the single label “AI.”
3. Vocabulary that matters	Which new terms affect price, risk, portability, and business value.
4. How AI is priced	Why seats, tokens, compute, storage, transactions, and services create different cost curves.
5. How to budget	How to calculate total cost of ownership rather than only the visible software bill.
6. Buy, build, partner?	Which capabilities should be purchased as commodities and which may justify proprietary investment.
7. Approving investment	A stage-gated framework for moving from idea to production capital.
8. Measuring return	Why cost per useful decision matters more than cost per token.
9. M&A case	How the framework changes a real information-intensive financial work-flow.
10. Procurement and contracts	What to ask vendors before the institution becomes dependent on them.
11. Board governance	Which questions belong at board level and which belong to management.
12. Executive decision rule	A concise framework for allocating AI capital with discipline.

The recurring mental model

Business Problem → AI Capability → Operating Architecture → Economics

↓

Controls + Ownership → Investment Decision

Do not start with the model. Start with the business problem. Then identify the minimum AI capability required, the architecture needed to operate it, the full cost of that architecture, and the controls necessary to make the result dependable.

Abstract

Artificial intelligence is difficult to buy because “AI” is not a single product category. A company may purchase a software application that happens to contain AI, a subscription to a general-purpose model, access to model inference through an API, cloud compute, a retrieval and knowledge layer, a custom agentic workflow, implementation services, or a vertically integrated private system. Each object has a different economic model, risk profile, degree of lock-in, and budgeting logic. This guide provides a non-technical but rigorous framework for CEOs, entrepreneurs, financial practitioners, and boards. It explains the minimum vocabulary necessary to evaluate vendors; decomposes total cost of ownership; distinguishes buy, build, and partner decisions; proposes a stage-gated investment process; and shows why the relevant unit of economic analysis is not the cost of a token or software seat but the *cost per useful, controlled business outcome*. A synthetic M&A screening case illustrates how architecture can change both cost and decision velocity. The central conclusion is that management should not approve “AI spending” as a generic technology category. It should approve specific business capabilities with explicit architectures, measurable economics, accountable owners, and defined control boundaries.

The foundational proposition

You are rarely buying “AI.”
You are buying a *business capability* built from several AI and non-AI layers.

What You Should Be Able to Do After Reading

You should be able to look at an AI proposal and identify what is actually being purchased; understand the major drivers of variable and fixed cost; distinguish a software subscription from a model service, an agent, a retrieval system, or a private architecture; decide when to buy, build, or partner; construct a practical AI budget; and ask the questions necessary to approve or reject an AI investment without becoming an AI engineer.

1. Why AI Feels Different from Technology You Have Bought Before

A CEO who has managed technology for many years already possesses most of the right instincts: insist on a business case, understand implementation costs, demand accountability, protect data, avoid vendor lock-in, and measure outcomes. What has changed is that the object being purchased has become less visible. Traditional enterprise software usually exposed a relatively clear boundary. The firm bought a license or a cloud service that performed a defined set of functions. AI systems are more fluid. Their behavior depends on models that may change, data that may be retrieved at runtime, usage that may fluctuate dramatically, and probabilistic outputs that may require verification.

The first source of confusion is therefore semantic. “We are buying AI” can describe transactions that are economically and operationally very different. A company buying an AI-enabled accounting package is still primarily buying enterprise software. A company calling a frontier model through an API is buying metered intelligence. A company deploying a private retrieval system is buying a knowledge architecture. A company building agents that can use tools and execute tasks is buying a new form of workflow automation. A company training its own model or reserving large GPU clusters is making something closer to an infrastructure commitment.

The second difference is that AI costs can move with usage in ways that are unfamiliar to executives accustomed to fixed software licenses. More employees, longer prompts, more documents, more reasoning, more model calls, or more autonomous steps may all increase variable cost. A proof of concept can therefore look inexpensive while the production system becomes costly at scale. Conversely, a well-designed architecture can reduce the amount of expensive model intelligence purchased for each unit of work.

The third difference is that AI changes faster than most enterprise technology categories. The model that appears indispensable today may become cheaper, weaker, stronger, or replaceable within a relatively short procurement cycle. This makes portability, modularity, and contractual exit options more strategically important than they were in many prior software categories.

Executive Principle 1: Preserve Your Existing Procurement Discipline

AI does not invalidate the disciplines that executives learned from decades of technology investment. It makes them more important. The new requirement is to understand the layers well enough to know *where* the cost, risk, and strategic dependency actually sit.

2. What Are You Actually Buying? The Seven-Layer View

The most useful way to remove ambiguity is to stop treating AI as a product and instead view it as a stack. A business outcome sits at the top. Several technical and organizational layers sit underneath. Different vendors may provide one layer, several layers, or the entire stack. A proposal becomes easier to evaluate once management can say precisely which layers are included and which remain the company's responsibility.

Layer	What It Is	What You May Be Buying	Executive Question
1. Outcome	The business result.	Faster underwriting, better customer service, lower fraud losses, more productive research, automated operations.	What measurable business result justifies the investment?
2. Workflow	The application or process through which people use AI.	Copilot, research workstation, call-center assistant, underwriting workflow, agentic process.	Does it fit the way work is actually done?
3. Intelligence	The model or models that interpret, generate, classify, predict, or reason.	Model subscription, API access, hosted open model, specialist model.	Are we dependent on one model, or can models be changed?
4. Knowledge	The mechanism that brings company information to the model.	Search, embeddings, vector database, RAG, knowledge graph, document pipelines.	How does the system know our facts rather than merely general facts?
5. Data & Integration	Connections to systems of record and enterprise data.	ERP/CRM connectors, APIs, data cleaning, identity and permissions.	What must be integrated before the AI can perform useful work?
6. Infrastructure	The compute, storage, networking, and hosting environment.	Cloud services, GPU capacity, private hosting, observability.	Who bears capacity, performance, and scaling cost?
7. Control	The governance and safety layer.	Evaluation, logging, access control, cyber controls, human approval, audit trails.	How do we know when the system is wrong, unsafe, or outside policy?

This seven-layer view immediately changes the procurement conversation. A vendor may offer an impressive model but no integration into the company's data. Another may offer a polished application but create a deep dependency on a proprietary knowledge store. A third may sell a complete managed solution with attractive simplicity but limited portability. None of these choices is intrinsically wrong. The mistake is approving the investment without understanding which layer creates the value and which layer creates the dependency.

A practical example

Suppose a bank purchases an "AI credit analyst." The visible product may be a conversational interface. Underneath, the vendor may be providing model inference, document ingestion, embeddings, a vector database, connectors to financial data, workflow logic, monitoring, and cloud hosting. The bank is not buying a chatbot. It is buying an integrated analytical production system. The budget and due diligence should reflect the full system, not the visible interface.

3. The New Vocabulary: Only the Terms That Change the Buying Decision

Executives do not need the full language of machine learning. They do need the terms that change economics, risk, or portability. The objective is not technical fluency for its own sake. The objective

is procurement fluency.

Term	Plain-English Meaning	Why It Matters When Buying
Foundation model / LLM	A general-purpose model capable of understanding and generating language and, increasingly, other modalities.	The model may be only one component of the product. Ask whether the vendor owns it, rents it, or can substitute another model.
Inference	Running a trained model to produce an answer or decision.	Inference is often a recurring usage cost. More calls, longer contexts, and more reasoning can increase OPEX.
Token	A small unit into which text is divided for model processing.	Some providers price model usage by input and output tokens. Token price is therefore one component of variable cost, not the business case itself.
Context window	The amount of information the model can consider in one interaction.	Large context can be useful but may increase cost and latency. A large advertised context window does not eliminate the need for information architecture.
Embedding	A numerical representation of meaning used to search semantically similar information.	Embeddings often power enterprise search and RAG. They can reduce how much raw content must be sent to an expensive model.
Vector database	A database optimized to store and search embeddings.	It is part of the enterprise knowledge layer. Ask about portability, security, indexing cost, and whether the vendor controls this layer.
RAG	Retrieval-Augmented Generation: retrieving relevant company information before asking the model to answer.	RAG can improve grounding, reduce unnecessary context, and create citations. It is frequently central to enterprise AI architecture.
Fine-tuning	Additional training that adapts a model's behavior to a narrower task or style.	It may improve specialized performance but creates development, evaluation, and maintenance obligations. Do not fine-tune merely because it sounds proprietary.
Agent	A model-driven software component that can choose steps and use tools to pursue a task.	Agents can create major productivity gains but also multiply model calls, permissions, and control requirements. Their cost is often per workflow, not per single prompt.
Tool use	Allowing the model or agent to call databases, applications, calculators, search systems, or transaction systems.	This is where AI moves from answering to acting. Permissions, authentication, audit logs, and error boundaries become critical.
Knowledge graph	A structured representation of entities and relationships in the firm's knowledge.	Useful where relationships, provenance, identity, and institutional knowledge matter. It may be a strategic data asset rather than a commodity software feature.
Evaluation / guardrails	Tests and controls that measure whether the system behaves acceptably.	These are part of production cost. A system without evaluation is difficult to govern, compare across model upgrades, or defend to a board.

Term	Plain-English Meaning	Why It Matters When Buying
GPU / accelerator	Specialized hardware used to train and run modern AI models efficiently.	Relevant if the company is buying dedicated capacity or private infrastructure. Often irrelevant if the vendor fully abstracts the infrastructure.

Executive Principle 2: Ask What the Term Changes

Whenever a technical term appears in a proposal, ask: *Does this change price, performance, data ownership, operational dependency, or risk?* If not, it may be implementation detail. If it does, it belongs in the investment discussion.

4. How AI Is Sold—and Why the Pricing Unit Matters

Traditional enterprise software trained executives to think in licenses, seats, projects, and maintenance contracts. Those models still exist, but AI introduces several additional pricing units. The unit chosen by the vendor determines which party bears the risk of usage growth.

Pricing Model	Typical Purchase	Economic Character	What to Watch
Per seat	Copilots and AI-enabled SaaS	Predictable fixed OPEX per user	Whether adoption is high enough to justify enterprise-wide licenses.
Per token	Frontier model APIs	Variable OPEX linked to text processed and generated	Context length, output length, model tier, reasoning intensity, and hidden repeated calls.
Per API call / transaction	Classification, extraction, workflow services	Variable cost linked to activity volume	A single business transaction may trigger multiple internal model calls.
Compute time / GPU	Hosted models, private deployments, training	Capacity-style OPEX or CAPEX-like commitment	Utilization risk: reserved capacity can be expensive if demand is uncertain.
Storage / vector volume	RAG, enterprise search, knowledge systems	Persistent data-layer cost	Re-indexing, document churn, replication, and vendor egress/portability.
Enterprise license	Integrated AI platforms	More predictable but may bundle many layers	Whether model usage caps or overage charges sit behind the headline fee.
Professional services	Integration, customization, controls, change management	Project cost, often front-loaded	The production system may require much more than the initial pilot integration.
Outcome / success fee	Selected vertical applications	Cost linked to business result	Definition of the outcome, measurement rights, and whether incentives encourage excessive automation.

The key lesson is that the visible price is often not the economic unit that management should optimize. A low token price can still produce an expensive workflow if an agent makes twenty calls per task. A high per-seat license can be attractive if it replaces several systems and materially changes throughput. A private model can appear expensive on infrastructure but become economical at very high and stable utilization. The CEO therefore needs the internal unit of work: one credit memorandum, one customer interaction, one contract reviewed, one product design cycle, one transaction screened, or one claim processed.

The right denominator

Do not optimize cost per token.
Optimize cost per useful, controlled business outcome.

5. How to Budget AI: From Vendor Price to Total Cost of Ownership

The most common budgeting error is to treat the model bill or software subscription as the AI budget. Production AI is an operating system of data, models, workflows, controls, people, and infrastructure. A useful total-cost formulation is:

$$\text{AI TCO} = C_{\text{software/model}} + C_{\text{data}} + C_{\text{integration}} + C_{\text{knowledge}} \\ + C_{\text{infrastructure}} + C_{\text{control}} + C_{\text{human}} + C_{\text{change}} + C_{\text{maintenance}}.$$

Each term has a different behavior. Model or software cost may scale with usage. Data and integration costs are often front-loaded but recur when systems change. Knowledge architecture requires ingestion, indexing, permissions, and refresh. Control costs include evaluation, monitoring, logging, cyber controls, and human escalation. Human costs include domain experts who validate outputs, redesign processes, and own exceptions. Change-management cost includes training, adoption, policy changes, and new performance incentives. Maintenance includes model upgrades, prompt/workflow revisions, re-testing, and vendor changes.

A disciplined budget therefore separates **build cost**, **run cost**, and **scale cost**. Build cost answers: what is required to create the production capability? Run cost answers: what does it cost at today's usage? Scale cost answers: what happens if adoption, transaction volume, context size, or autonomous activity increases fivefold? The third question is where many attractive pilots fail.

A practical AI budget should contain at least five lines

- 1. Business application and licenses.** The visible software or model access.
- 2. Data, integration, and knowledge architecture.** Connectors, retrieval, document pipelines, permissions, and data quality.
- 3. Variable intelligence cost.** Tokens, model calls, compute, agent steps, and high-cost reasoning tiers.
- 4. Control and resilience.** Evaluation, monitoring, cybersecurity, audit trails, human approval, and fallback procedures.
- 5. Organizational adoption.** Process redesign, training, operating ownership, and change management.

5.1 Budget by Workload, Not by “AI Department”

The CFO should resist a single undifferentiated AI budget. The more informative approach is to budget by business workload. Customer service, software development, legal review, credit analysis, marketing, research, and operations may use different models and have very different unit economics. A portfolio view allows management to compare one use case against another and to identify where shared platform investments—identity, retrieval, evaluation, data connectors, or model gateways—can be reused.

5.2 Treat the Platform as Shared Infrastructure

Certain AI costs should not be charged as if every use case were starting from zero. A secure model gateway, common evaluation framework, enterprise retrieval layer, permission system, and logging environment may support many applications. This resembles the evolution of cloud and data platforms: the firm benefits when foundational capabilities are standardized, while business units remain accountable for the economics of their specific use cases.

A pilot may involve fifty users, a small document set, manual quality checks, and forgiving latency. Production may involve five thousand users, continuous ingestion, automated actions, strict uptime, audit logging, security reviews, and a dramatically larger model-call volume. Always require a **production-scale cost model** before declaring a pilot economically successful.

6. Buy, Build, or Partner?

The classic technology decision remains valid, but AI changes the boundary. The firm should buy what is becoming commodity, build where proprietary information or workflow creates durable advantage, and partner where capabilities are strategically important but changing too quickly for the firm to internalize efficiently.

Choice	Best Fit	Advantages	Principal Risk
Buy	Commodity or near-commodity workflows; mature SaaS categories.	Speed, lower internal engineering burden, clearer vendor accountability.	Lock-in, bundled economics, limited differentiation, data portability.
Build	Proprietary workflow, unique data, differentiated decision process, strategic intellectual capital.	Control, differentiation, portability, ability to optimize architecture.	Engineering complexity, talent, maintenance, model churn, governance burden.
Partner	High-value use case where internal capability is incomplete or market architecture is still moving.	Shared expertise, faster learning, optionality, staged transfer of capability.	Dependency can become permanent if knowledge transfer and exit are not designed.

A useful test is to ask where the institution expects its competitive advantage to reside. If the advantage comes from the same general-purpose office productivity available to every competitor, buying is usually rational. If the advantage comes from the company’s proprietary underwriting logic, institutional knowledge, customer network, transaction history, or specialized operating process, then the architecture around the model may deserve proprietary investment even if the underlying model is purchased from a third party.

This leads to one of the most important strategic insights in AI procurement: **owning the model is not the same as owning the advantage**. A company can rent a frontier model and still own a highly differentiated knowledge graph, retrieval system, workflow, toolchain, evaluation framework, and institutional data asset. Conversely, a company can host an open model privately and still create little strategic value if the workflow around it is generic.

Executive Principle 3: Build the Differentiation, Buy the Commodity

The question is not “Should we build our own AI?” The better question is “Which layer of the AI stack contains our strategic advantage, and which layers should remain interchangeable commodities?”

7. How to Approve an AI Investment: A Stage-Gated Capital Process

AI investments should not move directly from enthusiasm to enterprise rollout. A disciplined process reduces both financial waste and governance risk. The stages below are deliberately familiar to executives because they resemble ordinary capital allocation. What changes is the evidence required at each stage.

Gate	Question	Evidence Required	Approval Logic
0. Problem	Is there a valuable problem?	Baseline cost, cycle time, error rate, revenue opportunity, risk exposure.	No AI discussion until the business problem is measurable.
1. Feasibility	Can AI improve it?	Small prototype; quality benchmark; data availability; workflow fit.	Fund learning, not scale.
2. Architecture	What must be built around the model?	Data map, integrations, retrieval, controls, human escalation, vendor dependencies.	Approve only when the production stack is visible.
3. Economics	Does the business case survive real usage?	Production TCO, unit cost, adoption scenario, sensitivity analysis, fallback cost.	Compare against current process and non-AI alternatives.
4. Controlled Pilot	Does it work with real users and controls?	Accuracy, override rate, failure taxonomy, user adoption, control performance.	Scale only if benefits persist under operational conditions.
5. Scale	Is the organization ready?	Capacity, resilience, cybersecurity, support, governance owner, vendor terms, exit plan.	Release scale capital with explicit operating accountability.
6. Review	Is it still worth funding?	Realized ROI, model drift, vendor changes, cost trend, incident data, user behavior.	Continue, redesign, re-source, or terminate.

This process protects against a recurring organizational error: celebrating technical feasibility as if it were an investment case. A model producing an impressive demo proves that a capability may be possible. It does not prove that the capability is reliable enough, integrated enough, controlled enough, adopted enough, or economical enough to deserve production capital.

The CEO’s investment sequence

Problem → Baseline → Prototype → Architecture → TCO → Controlled Pilot → Scale

8. How to Measure Return: The Unit Economics of AI

AI ROI should be evaluated at the level of the business process. The numerator is the value created or losses avoided. The denominator is the full incremental cost of producing that result with acceptable control. This can be expressed as:

$$\text{AI ROI} = \frac{\Delta \text{Revenue} + \text{Labor Capacity Released} + \text{Losses Avoided} + \text{Capital Efficiency Gained} - \text{AI TCO}}{\text{AI TCO}}.$$

The equation is intentionally broader than headcount reduction. Many high-value AI investments do not eliminate jobs. They increase throughput, reduce cycle time, allow professionals to examine more alternatives, improve conversion, reduce errors, detect risk earlier, or increase the quality of decisions. In professional services and financial businesses, the most valuable benefit may be that scarce senior judgment is applied to more cases rather than consumed by information gathering.

Management should therefore establish a baseline before implementation. How long does the process take today? What does it cost? What proportion of cases require rework? What is the error or exception rate? How much revenue is lost because decisions take too long? How many opportunities are not evaluated because analytical capacity is scarce? Without this baseline, almost any AI demonstration can appear productive.

8.1 The Four Sources of Economic Value

A useful executive taxonomy separates four sources of value. **Efficiency value** reduces the cost of existing work. **Throughput value** allows more work to be completed with the same resources. **Decision-quality value** improves risk-adjusted outcomes, conversion, pricing, or capital allocation. **Strategic-option value** creates a capability that enables future products or operating models. These categories should not be conflated because they require different evidence.

The economic question to ask every sponsor

“If this system is successful, what unit of business activity becomes materially better—and by how much?”

The answer should be something measurable: cost per claim, hours per diligence, conversion per lead, fraud loss per account, software release cycle, time to credit decision, revenue per relationship manager, or investment opportunities evaluated per analyst. “Employees will use AI” is not an investment thesis.

9. A Practical Case: What Is the CEO Buying in an M&A Screening System?

Consider a synthetic investment bank screening 500 potential acquisition targets for a corporate client. The corpus includes regulatory filings, earnings transcripts, credit agreements, and patents. The banker wants to identify financially attractive targets with appropriate leverage, proprietary industrial technology, and acceptable customer concentration. A naive procurement description might say: “We need an LLM to analyze the targets.” That description is too vague to budget or govern.

The actual production capability contains several purchases. The bank needs structured financial data and document ingestion. It needs an enterprise knowledge layer that can retrieve relevant passages. It needs model inference for extraction and synthesis. It needs an interface integrated into the banker’s workflow. It needs permissions, citations, logging, and human review. It may need several models rather than one: inexpensive models for classification and extraction, stronger reasoning models for difficult comparisons, and deterministic tools for financial calculations.

This is where architecture changes economics. In the synthetic case, asking a frontier model to repeatedly process the full universe is modeled at \$18,450 per sweep and 18.5 hours of wall-clock time. A four-tier retrieval-and-reasoning architecture is modeled at \$215 and 42 seconds per sweep. Across 30 iterations, the modeled cost changes from \$553,500 to \$6,450. These figures are illustrative rather than empirical benchmarks, but they reveal a crucial procurement principle: **the model price is not the system economics**. The architecture determines how much expensive intelligence must be purchased for each question.

Component	What the Bank Is Buying	Why It Matters
Structured screening	Financial data, deterministic rules, metadata filters	Removes companies that obviously fail quantitative criteria.
Retrieval layer	Embeddings, lexical search, vector index, document pipeline	Finds the small number of passages relevant to each question.
Re-ranking	More precise relevance model	Distinguishes merely related text from probative evidence.
Bounded inference	LLM/SLM extraction and reasoning	Uses model intelligence on a much smaller evidence set.
Workflow	Screening interface, structured output, analyst controls	Converts technical capability into a usable banking process.
Governance	Citations, logging, approvals, evaluation	Makes outputs reviewable by bankers, risk, compliance, and clients.

The CEO or board does not need to choose the embedding dimension or vector-database index. It does need to know that the proposed system avoids paying frontier-model economics for tasks that cheaper tools can solve; that critical conclusions can be traced back to evidence; that the bank owns or can export its data and knowledge layer; and that the system remains useful if the preferred model vendor changes.

Executive Principle 4: Architecture Is an Economic Decision

Retrieval, model routing, caching, smaller specialist models, deterministic tools, and human review are not merely engineering details. They determine the marginal cost, reliability, and substitutability of the system. Those are capital-allocation issues.

10. Procurement and Contracts: What to Ask Before You Become Dependent

AI procurement should examine more than functionality. Because the underlying models and economics can change rapidly, the contract should preserve the institution’s ability to understand, control, and exit the relationship. The precise legal language belongs to counsel, but executives should know which commercial questions matter.

- What exactly is included?** Model access, retrieval, storage, connectors, monitoring, implementation, support, and future upgrades should be separated conceptually even if bundled commercially.
- How is usage measured?** Ask for the pricing unit, included volume, overage rates, model-routing rules, and whether hidden internal calls can multiply usage.

3. **What happens to our data?** Determine where data is stored, how long it is retained, who can access it, whether it is used to improve vendor models, and how deletion is verified.
4. **Can we export our knowledge layer?** The firm's document indexes, embeddings, prompts, evaluations, workflow logic, and structured outputs may become valuable institutional assets. Portability should be considered early.
5. **Can the underlying model be changed?** A product that is inseparable from one model provider may create more dependency than a modular architecture with a model gateway.
6. **How are model changes communicated?** A vendor can improve or replace a model and thereby alter system behavior. Material changes should trigger re-evaluation before critical workflows silently drift.
7. **What evidence and logs are preserved?** For high-stakes uses, management should be able to reconstruct which data, model, tools, and approvals produced an important output.
8. **What are the service and resilience commitments?** Availability, latency, incident response, business continuity, and fallback modes matter when AI becomes part of operations.
9. **What is the exit plan?** Management should understand migration cost before dependency becomes strategic.

The easiest vendor to deploy may also be the hardest vendor to leave if it owns the workflow, knowledge store, model access, evaluation data, and user interface simultaneously. Bundling can be economically rational, but dependency should be a conscious investment choice rather than an accidental technical consequence.

11. What the Board Should Govern—and What It Should Not

Boards do not need to supervise prompts, model benchmarks, embedding algorithms, or GPU utilization. Those are management and technical responsibilities. The board's role is to ensure that AI investment fits strategy, capital discipline, risk appetite, and accountability. The board should be able to understand where AI is becoming material to revenue, cost, customer treatment, capital allocation, operational resilience, or regulated decision-making.

Board Question	Why It Matters	Evidence Management Should Provide
Where are we investing?	Prevents a diffuse portfolio of experiments with no strategic concentration.	Portfolio by use case, business owner, stage, spend, and expected value.
What is actually proprietary?	Distinguishes rented model capability from institutional advantage.	Map of owned data, workflows, knowledge assets, controls, and vendor dependencies.
What is the production TCO?	Makes scale economics visible before adoption accelerates.	Fixed/variable cost model and usage sensitivities.
Where can AI act autonomously?	Autonomous tool use changes operational and control risk.	Permission model, approval thresholds, transaction limits, and human escalation.
Can important outputs be reconstructed?	Necessary for accountability, audit, and incident response.	Logging, provenance, versioning, and evidence-retention design.
How dependent are we on vendors?	Concentration can become strategic and financial risk.	Exit plan, model portability, data portability, alternative suppliers.
Is value being realized?	Prevents AI spending from becoming an innovation entitlement.	Realized ROI, adoption, quality, incidents, and discontinued projects.

The CEO and executive team, meanwhile, should own the operating model. Every production AI capability needs a business owner who is accountable for value, a technical owner accountable for architecture and reliability, and a control owner accountable for risk boundaries. In many institutions the CFO should also play a stronger role than in early AI experimentation because variable inference and cloud economics can create meaningful operating leverage—positive or negative.

The board-level test

Before approving material AI scale-up, the board should be able to answer seven questions in plain language: *What problem are we solving? What are we actually buying? What will it cost at scale? What do we own? What are we dependent on? What can the system do without a human? How will we know whether it is creating value?*

12. The Executive Decision Rule: How to Allocate AI Capital

The most useful conclusion is not that every company should spend aggressively on AI, nor that every firm should wait for the technology to mature. The rational position is more disciplined. AI should be treated as a portfolio of business-capability investments whose technical ingredients are changing rapidly. That argues for strong experimentation, but also for modular architecture, explicit unit economics, and staged capital commitments.

The executive decision can be summarized in six steps. First, define the business problem without reference to AI. Second, identify the minimum AI capability required to improve the problem. Third, decompose the proposed solution into the seven layers so that management knows what is being purchased and what remains internal. Fourth, model total cost at production scale rather than pilot scale. Fifth, decide which layers should be bought, built, or supplied through a partner based on strategic differentiation. Sixth, release scale capital only after the system demonstrates value under real controls and real usage.

The six-question CEO checklist

1. **What business outcome am I funding?**
2. **What exactly am I buying—application, model, knowledge layer, infrastructure, services, or all of them?**
3. **What is the full cost to build, run, and scale it?**
4. **Which part should we own because it differentiates us?**
5. **What must be true for the system to be trusted and governed?**
6. **What evidence will cause us to scale, redesign, re-source, or stop?**

For the non-specialist CEO, this is the central reassurance: the investment decision does not require mastery of transformer mathematics or GPU engineering. It requires a clear mental model of what the proposed capability is made of, how its costs behave as usage grows, where strategic advantage can reside, which dependencies are being accepted, and which risks become material when software begins to interpret, recommend, and act. Once those questions are separated, AI becomes less mysterious. It becomes what executives have always managed: a capital-allocation problem under uncertainty, with technology as the instrument and business value as the objective. The final checklist converts that mental model into a repeatable approval discipline that can be used in management meetings, capital committees, procurement reviews, and board discussions.

Do not approve an AI budget.

**Approve a portfolio of business capabilities,
each with an architecture, an owner, a cost curve, a control
framework, and a measurable return.**

13. Executive AI Investment Checklist: Before You Approve the Spend

This checklist is designed to be used at the moment an AI proposal moves from discussion to commitment. It is intentionally written in business language. A CEO, CFO, investment committee, or board member should be able to work through it without knowing how to train a model or configure a vector database. The purpose is not to create bureaucracy. It is to prevent the organization from approving an attractive demonstration without understanding the full capability, the economics at scale, the ownership structure, or the control obligations that come with it.

How to use the checklist

For each item, management should be able to answer **Yes**, **Not Yet**, or **Not Applicable**, and point to the evidence supporting the answer. A material “Not Yet” is not necessarily a reason to cancel the initiative; it is a reason to avoid pretending that the investment is ready for the next gate.

A. Business Case: Are We Funding a Real Problem?

- ☐ Can we describe the business problem in one sentence without using the words “AI,” “LLM,” or “agent”?
- ☐ Do we know the current baseline: cost, cycle time, error rate, capacity constraint, revenue leakage, risk exposure, or customer friction?
- ☐ Is there a named business owner accountable for the outcome, not merely a technology sponsor?
- ☐ Have we defined the unit of value that should improve—for example, cost per claim, time per credit decision, revenue per adviser, contracts reviewed per lawyer, or opportunities screened per analyst?
- ☐ Have we considered whether ordinary software, process redesign, analytics, or automation could solve the problem more cheaply?
- ☐ Is the expected value large enough to justify not only a pilot, but the organizational attention required to operate the capability in production?

B. The Purchase: Do We Know What We Are Actually Buying?

- ☐ Can management identify which of the seven layers are included: outcome, workflow, intelligence, knowledge, data/integration, infrastructure, and control?
- ☐ Are we buying an application, model access, a knowledge system, an agentic workflow, infrastructure, implementation services, or a bundled combination?
- ☐ Do we know which components are supplied by the vendor and which remain our responsibility?
- ☐ Do we know which model or models sit underneath the product, and whether they can be changed without replacing the whole system?
- ☐ Do we understand how the system will obtain our institutional facts and documents rather than relying only on general model knowledge?

- ☐ If the product is described as an “agent,” do we know exactly what tools, systems, data, or transactions it is allowed to access or execute?

C. Budget and Unit Economics: Do We Know What It Will Really Cost?

- ☐ Does the budget include more than the visible license or model invoice?
- ☐ Have we separated **build cost**, **run cost**, and **scale cost**?
- ☐ Does total cost of ownership include data preparation, integrations, knowledge architecture, infrastructure, evaluation, cybersecurity, human review, change management, support, and maintenance?
- ☐ Do we know the vendor’s pricing unit—seat, token, API call, transaction, compute hour, storage, enterprise license, or outcome?
- ☐ Do we know the internal economic unit that matters to us—the cost per useful, controlled business outcome?
- ☐ Have we modeled at least a base case and a high-adoption case, including longer contexts, more users, more agent steps, and higher reasoning intensity?
- ☐ Do we understand which costs are fixed, which are variable, and which could increase nonlinearly as usage grows?
- ☐ Have we budgeted for re-testing, model changes, vendor migrations, and continuing evaluation after launch?

D. Buy, Build, or Partner: Are We Owning the Right Things?

- ☐ Have we identified which layer, if any, contains genuine strategic differentiation for our firm?
- ☐ Are we buying commodity capabilities rather than rebuilding them unnecessarily?
- ☐ Are proprietary data, institutional knowledge, workflow logic, evaluations, and structured outputs owned or exportable by us where strategically important?
- ☐ If we are building internally, is the expected advantage large enough to justify engineering, maintenance, talent, and governance costs?
- ☐ If we are partnering, is knowledge transfer explicit so that partnership does not become accidental permanent dependence?
- ☐ Can we change model providers, infrastructure providers, or application vendors without recreating the entire capability from zero?

E. Trust, Control, and Accountability: Can We Operate It Safely?

- ☐ Is there a defined accuracy or quality threshold appropriate to the business decision being supported?
- ☐ Do we know the most important failure modes and how they will be detected?
- ☐ Can material outputs be traced to source evidence, data, model version, tools used, and approvals where appropriate?

- ☐ Is there a clear human-approval boundary for high-impact decisions or actions?
- ☐ Are permissions designed so the system can access only the data and tools required for its role?
- ☐ Have cybersecurity, privacy, confidentiality, regulatory, and record-retention requirements been incorporated into the production design?
- ☐ Is there a fallback process when the AI service is unavailable, uncertain, or outside policy?
- ☐ Is someone explicitly accountable for model/workflow performance after deployment, including drift, incidents, overrides, and exceptions?

F. Vendor and Contract: Can We Understand and Exit the Relationship?

- ☐ Does the contract make clear what is included, what is metered, and what triggers overage charges?
- ☐ Do we know whether our data, prompts, outputs, or documents may be retained or used to improve vendor systems?
- ☐ Are material model or product changes communicated so that critical workflows can be re-tested?
- ☐ Do service commitments address availability, latency, resilience, support, and incident response where operationally relevant?
- ☐ Can we export our data, knowledge artifacts, evaluations, logs, and other strategically important assets in usable form?
- ☐ Have we estimated the operational and financial cost of leaving the vendor before lock-in becomes material?

G. Pilot-to-Scale Decision: Have We Earned the Right to Scale?

- ☐ Was the pilot tested with real users, realistic data, and production-like controls rather than only a curated demonstration?
- ☐ Did the pilot improve the pre-defined business baseline rather than merely receive positive qualitative feedback?
- ☐ Do we know the override rate, exception rate, major failure categories, and level of human effort still required?
- ☐ Has adoption been measured rather than assumed?
- ☐ Does the business case still work using production-scale costs rather than pilot costs?
- ☐ Is the organization prepared to support the capability operationally after the project team leaves?
- ☐ Have we defined in advance the evidence that would cause us to scale, redesign, re-source, pause, or terminate the system?

The One-Page Approval Test

If time is limited, the executive team should at minimum be able to answer the following twelve questions before authorizing a material AI deployment.

#	Approval Question	Status
1	What measurable business outcome are we buying?	☐ Yes ☐ Not Yet
2	What is the current baseline against which success will be measured?	☐ Yes ☐ Not Yet
3	What exactly are we buying across the seven layers of the system?	☐ Yes ☐ Not Yet
4	What is the full build-run-scale cost, not just the vendor price?	☐ Yes ☐ Not Yet
5	What is the cost per useful, controlled business outcome at expected scale?	☐ Yes ☐ Not Yet
6	What should we own because it creates differentiation, and what should remain commodity?	☐ Yes ☐ Not Yet
7	Which vendor, model, data, or infrastructure dependencies are we accepting?	☐ Yes ☐ Not Yet
8	Can critical outputs be reconstructed and traced to evidence?	☐ Yes ☐ Not Yet
9	What may the system do without a human, and where is approval mandatory?	☐ Yes ☐ Not Yet
10	Who owns business value, technical reliability, and control performance after launch?	☐ Yes ☐ Not Yet
11	What is our exit or re-sourcing plan if economics, models, or vendors change?	☐ Yes ☐ Not Yet
12	What evidence will make us scale, redesign, re-source, or stop?	☐ Yes ☐ Not Yet

A material AI investment should normally **not move to scale** if any of the following remain unresolved: **(1)** no measurable business baseline; **(2)** no production-scale TCO; **(3)** no accountable business owner; **(4)** no defined control, provenance, or human-escalation model for material decisions; or **(5)** no understanding of vendor dependency and exit. A strong demo does not compensate for any of these omissions.

The final executive rule

Approve AI when you can explain the **business outcome,
the **architecture**, the **economics**, the **ownership**,
and the **control model** in plain language.**

If the proposal can only be defended through technical vocabulary, vendor enthusiasm, or fear of being left behind, the investment case is not yet complete.