SYSTEMS | AGENTS | AUTONOMY | KNOWLEDGE

# From Software Factory to Autonomous Financial Institution

What Uber's Agentic Architecture Teaches Financial Professionals About the Next Stage of Artificial Intelligence

**Deep-Dive Pedagogical Guide ▪ Systems, Agents, Knowledge, Cybersecurity, Self-Improvement & Cooperation**

**Alejandro Reynoso**

Applied AI Architecture / Financial Systems

**September 2026**

Pedagogical Paper

A rigorous but non-specialist guide to understanding how production AI evolves from models to agents, autonomy, institutional knowledge, governed execution, self-improvement, and multi-agent cooperation.

## Reader's Map

**THE FINANCIAL PROFESSIONAL'S QUESTION**

**"I understand models, markets, risk, and institutions.
What changes when AI stops merely answering and begins to act?"**

The purpose of this paper is not to turn the financial professional into a software engineer. It is to provide enough architectural understanding to recognize what an agentic system is made of, why autonomous agency changes the governance problem, how institutional knowledge becomes a strategic asset, and why self-improving populations of agents eventually raise game-theoretic questions.

| Section | The question it answers |
|---|---|
| Why Uber matters | Why a software-engineering case is relevant to finance. |
| Understanding the system | Why the model is only one layer of institutional AI. |
| The financial cognitive factory | How financial work can be understood as an information-production system. |
| How agents work | What changes when AI moves from answering to acting. |
| Skills and orchestration | How institutional procedures become executable and how agents divide work. |
| Autonomous agency | What changes when the machine no longer waits to be asked. |
| Institutional knowledge engineering | Why organizational memory becomes part of system intelligence. |
| Cybersecurity and harnessing | How permissions, identity, provenance, and control bound machine action. |
| Self-improvement | What happens when the system begins improving its own operating procedures. |
| Game theory and cooperation | Why interacting agents create a strategic and emergent layer. |
| The autonomous financial factory | How the architecture translates into finance. |
| Economics and advantage | Why architecture determines cost and may become the durable asset. |

**THE RECURRING MENTAL MODEL**

**Business Problem → AI System → Agentic Workflow → Autonomy**

↓

**Knowledge + Controls + Learning → Institutional Capability**

Do not start with the chatbot. Start with the business problem and the institutional process.

Then identify the intelligence, knowledge, tools, agents, authority boundaries, economics, and controls required to create a dependable capability.

## Abstract

Artificial intelligence is often introduced to financial professionals through a deceptively simple interface: a person asks a question, a large language model generates an answer, and productivity improves. That experience is useful, but it increasingly describes only the outermost layer of a deeper transformation. The more consequential development is the emergence of AI systems that operate inside institutional environments, retrieve knowledge, invoke tools, coordinate specialized capabilities, initiate work, monitor outcomes, correct failures, and increasingly improve the procedures through which they themselves operate.

Uber's 2026 description of its internal "Software Factory" provides an unusually concrete illustration of this transition [1]. The company reports that more than 70 percent of pull requests are attributed to local or cloud agents, that engineers have created more than 3,600 agent skills, and that these skills execute more than 30,000 times per day. More importantly, Uber describes managed agents that review code, repair failures, triage operational alerts, debug systems, complete software changes, and perform maintenance within a governed architecture.

For financial professionals, the significance of this case extends far beyond software engineering. A financial institution is also a cognitive production system. It transforms information, institutional knowledge, market observations, regulatory constraints, capital, models, and human judgment into decisions, transactions, portfolios, financing structures, and risk controls. This paper interprets Uber's architecture through six connected layers: understanding AI as a system; understanding how agents operate; autonomous agency; institutional knowledge engineering; cybersecurity and harnessing; and self-improving multi-agent systems characterized by cooperation, strategic interaction, and game-theoretic dynamics.

The central conclusion is that the next frontier of artificial intelligence in finance will not be defined by better chatbots. It will be defined by the institutional architecture through which models, agents, tools, knowledge, humans, controls, and increasingly autonomous forms of intelligence are assembled into a coherent operating system.

**THE FOUNDATIONAL PROPOSITION**

**The model is not the institution.**

Institutional intelligence emerges from the architecture through which models, knowledge, tools, agents, humans, and controls are connected.

## Why Uber's Software Factory Matters to Finance

At first reading, Uber's *Running a Software Factory Efficiently at Uber Scale* appears to be a sophisticated discussion of how to manage the rapidly increasing cost of artificial intelligence in software development [1]. The article does examine token cost, routing among models, context expansion, tool overhead, caching, benchmarking, and observability. Yet treating it merely as an efficiency paper would miss its deeper importance.

What Uber is really describing is an emerging operating architecture for machine agency. Models sit inside an environment of reusable skills, tool servers, agent orchestration, institutional context, evaluation systems, authentication, policy enforcement, trace collection, and managed autonomous workflows. Uber's own scale makes the case particularly revealing: more than 70 percent of pull requests are attributed to local or cloud agents, more than 3,600 agent skills have been created, and more than 30,000 skill executions occur per day [1]. The interesting fact is not simply that developers are receiving better code suggestions. The production process itself is increasingly being reorganized around machine actors.

For finance, the analogy is powerful because financial institutions are already systems that transform information into consequential decisions. A bank transforms data and judgment into credit, liquidity, pricing, and risk allocation. An asset manager transforms information and beliefs into portfolios. An investment bank transforms corporate, legal, market, and investor information into transactions. A securities exchange transforms orders, rules, surveillance, and market information into trading infrastructure. The common denominator is that these organizations are cognitive production systems.

The relevant question is therefore no longer simply whether an analyst can use an LLM to summarize an earnings call or whether an investment banker can use AI to accelerate a memorandum. Those are useful applications, but they still leave the human as the primary orchestrator of the process. The deeper transformation begins when artificial agents participate in the production architecture itself—retrieving context, invoking models, calling deterministic tools, dividing work, observing results, escalating exceptions, and operating continuously inside institutional boundaries.

**SECTION TAKEAWAY**

**Main idea.** Uber's Software Factory is important because it shows AI migrating from a user-facing productivity tool into the operating architecture of an institution.

**What this teaches us about AI systems.** The meaningful object of analysis is no longer only the model. It is the system of models, tools, skills, context, orchestration, controls, and feedback that turns intelligence into production.

**Main takeaway for financial modeling with AI.** Financial professionals should move from asking "Which model should we use?" toward asking "What complete analytical and operating system must surround the model for this financial capability to be reliable, economical, and governable?"

## A. Understanding the System: Why the Model Is No Longer the System

One of the most persistent misunderstandings in contemporary artificial intelligence is the tendency to equate AI with the large language model itself. The error is understandable. A user interacts with a frontier system through a conversational interface and therefore experiences the model as the product. Yet production AI increasingly resembles what researchers have called a compound AI system: several specialized components cooperate to generate an institutional outcome [10].

The familiar representation

$$\text{Prompt} \rightarrow \text{Model} \rightarrow \text{Response}$$

is useful for explaining a chatbot. It is inadequate for explaining institutional AI. A more realistic abstraction is

$$\mathcal{S} = \{M, C, T, K, A, O, G, E\},$$

where $M$ denotes models, $C$ context, $T$ tools, $K$ institutional knowledge, $A$ agents, $O$ orchestration, $G$ governance, and $E$ evaluation. The quality visible to the end user is therefore an emergent property of the architecture.

Uber makes this explicit in practice. The company does not simply ask which frontier model is strongest. It evaluates candidate models against real tasks using a common harness and compares quality, latency, cost, reliability, and task completion. A more capable model may plan and supervise while less expensive models perform narrower subtasks. Tool discovery is also managed architecturally because loading every possible capability into every agent would create unnecessary context and cost [1]. These are not merely engineering details. They are decisions about where intelligence should reside and how much expensive intelligence should be purchased for each unit of work.

The implication for finance is immediate. A portfolio-management system might combine a reasoning model with market data, internal research, current holdings, risk limits, transaction-cost models, portfolio optimization, regulatory rules, historical decisions, and human approval. A credit system might combine a model with borrower data, covenant databases, legal documents, rating histories, scenario engines, policy rules, and committee escalation. In each case, the model is important, but the model alone is not the capability.

The pedagogical shift is therefore from *model literacy* to *system literacy*. The financial professional does not need to become a machine-learning engineer, but does need to understand enough architecture to ask which information enters the system, which tools it can invoke, how outputs are evaluated, where authority resides, what happens when the system is uncertain, and how actions are reconstructed afterward.

**Executive Principle 1: Do Not Confuse the Model with the Capability**

A strong model can sit inside a weak system, and a well-designed system can often extract more value from a less expensive model. In institutional AI, architecture is part of intelligence.

**SECTION TAKEAWAY**

**Main idea.** Production AI is a system, not a model. Institutional capability emerges from the interaction of intelligence, context, tools, knowledge, orchestration, evaluation, and control.

**What this teaches us about AI systems.** The model is only one source of capability. Context quality, tool design, routing, memory, and governance can materially alter both performance and cost.

**Main takeaway for financial modeling with AI.** Financial modeling should evolve beyond selecting the "best model." The stronger design question is how to connect models with financial data, deterministic analytics, institutional knowledge, validation, optimization, risk limits, and human judgment.

## The Financial Institution as a Cognitive Production System

The factory metaphor is particularly useful for finance because financial institutions have always transformed information into decisions. Their raw materials are not steel or components. They are prices, contracts, balance sheets, mandates, regulatory rules, customer histories, market microstructure, risk exposures, and human judgments.

A commercial bank receives information about borrowers, deposits, collateral, interest rates, liquidity, regulation, and macroeconomic conditions. It transforms those inputs into loans, pricing, capital allocation, and risk positions. An asset manager receives market prices, research, alternative data, portfolio constraints, and investment beliefs and transforms them into allocations. An investment bank receives corporate, legal, valuation, financing, investor, and regulatory information and transforms it into transactions. The institution can therefore be represented as a cognitive production function:

$$\mathcal{Y} = F(\text{Information}, \text{Knowledge}, \text{Models}, \text{Judgment}, \text{Capital}, \text{Constraints}).$$

Historically, humans performed most of the cognitive transformation. Software then automated calculations, storage, routing, and predefined workflows. Generative AI added a new layer by allowing software to interpret and synthesize unstructured information. Agentic AI introduces something further: parts of the cognitive process can now be decomposed into persistent machine actors that use tools and operate over multiple steps.

The institutional progression is therefore better viewed as

$$\text{Human Institution} \rightarrow \text{Digitized Institution} \rightarrow \text{AI-Assisted Institution} \rightarrow \text{Agentic Institution} \rightarrow \text{Partially Autono}$$

The difference between the later stages is organizational rather than cosmetic. In the AI-assisted institution, the human remains the principal initiator and coordinator. In the agentic institution, the human can delegate a multi-step objective. In a partially autonomous institution, the system can monitor events, activate workflows, coordinate specialist agents, and execute bounded actions under policy.

This is why the Uber case is so useful pedagogically. It allows the financial professional to see AI not as an isolated software purchase but as a potential change in the production function of the institution itself.

**SECTION TAKEAWAY**

**Main idea.** A financial institution is already a cognitive factory: it converts information, institutional knowledge, constraints, and judgment into financial outcomes.

**What this teaches us about AI systems.** Agentic AI matters because it can participate directly in the cognitive production process rather than merely accelerating isolated tasks.

**Main takeaway for financial modeling with AI.** The strategic opportunity is not merely to automate spreadsheets or improve a forecast. It is to redesign analytical workflows so that data, models, knowledge, agents, optimization, and human judgment operate as one production system.

## B. How Agents Work: The Difference Between Answering and Acting

The term *agent* is often used too loosely. A language model that produces a response is not automatically an agent. Agency begins to emerge when the model is embedded in a loop that allows it to interpret a state, pursue an objective, select an action, observe the consequence, and continue.

A simplified representation is

$$s_t \to \pi \to a_t \to o_{t+1},$$

where $s_t$ is the state available to the agent, $\pi$ its decision or reasoning policy, $a_t$ the action selected, and $o_{t+1}$ the resulting observation. The important feature is continuation. A conventional chatbot answers. An agent proceeds.

It may search for information, query a database, invoke a pricing engine, inspect a document, call another model, ask another agent to investigate a subproblem, compare outputs, discover a contradiction, revise its plan, retry a failed tool call, or escalate to a human. The pattern is therefore

$$\text{Observe} \to \text{Reason} \to \text{Act} \to \text{Evaluate} \to \text{Continue}.$$

This structure is closely related to the reasoning-and-action architecture formalized in ReAct [2]. What makes the architecture consequential for finance is that financial work is naturally multi-step. A credit recommendation is not one prediction. An investment thesis is not one inference. A merger screen is not one question. Each process involves retrieval, calculation, comparison, judgment, documentation, and review.

Consider a portfolio agent asked whether the institution should increase exposure to a company. It might retrieve the latest financial statements, identify changes in guidance, invoke a valuation model, inspect credit spreads, compare the position against portfolio factor exposures, retrieve the previous investment committee thesis, run scenario analysis, and finally produce a recommendation with explicit evidence. If uncertainty remains high, it may escalate.

The important difference is that the AI is no longer only generating text. It is participating in a workflow.

**SECTION TAKEAWAY**

**Main idea.** An agent is defined by its ability to remain inside a task loop, use tools, observe results, and continue until an objective is completed or escalation is required.

**What this teaches us about AI systems.** Agency introduces persistence, tool use, feedback, and multi-step execution. The system becomes an actor in an environment rather than merely a generator of responses.

**Main takeaway for financial modeling with AI.** A financial model becomes much more powerful when embedded in an agent that can obtain data, run calculations, compare hypotheses, validate outputs, query institutional knowledge, and determine when human review is necessary.

## Skills and Orchestration: From Procedure to Executable Capability

Uber's report of more than 3,600 agent skills is one of the most revealing facts in the case [1]. A skill should not be understood merely as a shortcut. It is a codified capability that converts general-purpose intelligence into repeatable institutional action.

A useful abstraction is

$$\text{Skill} = \text{Instructions} + \text{Context} + \text{Tools} + \text{Procedure} + \text{Constraints}.$$

Financial institutions already contain thousands of analogous capabilities. Credit underwriting is a skill. Comparable-company analysis is a skill. Covenant review is a skill. Liquidity stress testing is a skill. Portfolio reconciliation is a skill. Fraud investigation is a skill. Historically, these capabilities have lived in people, manuals, spreadsheets, code, workflow systems, and institutional custom. Agent architectures make at least part of them machine-executable.

The progression is therefore

$$\text{Tacit Practice} \rightarrow \text{Documented Procedure} \rightarrow \text{Machine-Readable Skill} \rightarrow \text{Agent-Executable Capability}.$$

This is potentially a new form of institutional capital. A senior credit officer may possess decades of judgment about which questions matter, which covenants deserve scrutiny, and which borrower behaviors signal hidden weakness. Not all of that expertise can be formalized, but part of the procedure can increasingly be captured, linked to evidence, connected to tools, and reused by agents.

Once many skills exist, orchestration becomes the next problem. A capable primary agent does not need to perform every subtask itself. It can decompose the objective and allocate work to specialist agents:

$$A_0 \rightarrow \{A_1, A_2, \ldots, A_n\} \rightarrow A_0.$$

The orchestrator behaves less like a worker and more like a manager: it defines the problem, divides the work, selects the appropriate specialists, evaluates intermediate outputs, and synthesizes the result. This logic appears in multi-agent frameworks such as AutoGen [3] and is especially intuitive in finance because professional organizations already rely heavily on specialization.

A transaction, for example, may require legal, valuation, tax, credit, market, regulatory, and structuring perspectives. A computational architecture can mirror this division of labor without pretending that one monolithic model should perform every task equally well.

**Executive Principle 2: Treat Skills as Institutional Capital**

The long-run strategic asset may not be a single prompt or model. It may be the accumulated library of governed, reusable, institution-specific skills that encode how the organization actually performs valuable work.

**SECTION TAKEAWAY**

**Main idea.** Skills convert institutional procedures into reusable agent capabilities, while orchestration allows specialized agents to divide complex work.

**What this teaches us about AI systems.** General intelligence becomes much more useful when decomposed into repeatable procedures and coordinated through hierarchy, specialization, and evaluation.

**Main takeaway for financial modeling with AI.** Financial firms should identify which analytical practices can become reusable skills and which complex processes can be decomposed into specialist-agent roles resembling computational deal teams, credit committees, or investment committees.

## C. Autonomous Agency: When the Machine Stops Waiting to Be Asked

The transition from agentic workflows to autonomous agency occurs when the system no longer requires a person to initiate every task. Uber explicitly describes managed agents that perform code review, repair broken continuous-integration pipelines, triage alerts, debug defects, and conduct maintenance [1]. The architectural shift can be represented simply:

$$\text{Human Request} \rightarrow \text{Agent Action}$$

becomes

$$\text{Event} \rightarrow \text{Agent Action.}$$

This change is more important than it may initially appear. When a system waits for a human prompt, the human remains the initiating intelligence. When the environment itself can activate an agent, limited initiative has moved into the machine.

Finance contains many naturally event-driven processes. A risk system today may answer the question, "What happens to the portfolio if rates rise by 100 basis points?" An autonomous risk architecture would observe the rate movement, identify the exposed books, invoke scenario engines, inspect liquidity, examine counterparties, compare results against limits, generate possible mitigation actions, and escalate only when predefined thresholds are crossed.

The human no longer needs to formulate the first question. The system has acquired bounded initiative.

Autonomy should nevertheless be understood as a spectrum rather than a binary state. The relevant design problem is not to maximize autonomy but to optimize useful autonomy subject to constraints:

$$\max \ \text{Useful Autonomy}$$

subject to

$$\text{Risk Limits}, \quad \text{Permissions}, \quad \text{Legal Constraints}, \quad \text{Auditability}, \quad \text{Human Escalation.}$$

A trading agent may be authorized to rebalance inside narrow bands but not initiate a new strategic exposure. A treasury agent may optimize liquidity within an approved mandate but not alter the funding strategy. A surveillance agent may open an investigation automatically but not impose a sanction. The architecture therefore resembles familiar delegation systems in finance, except the delegated actor is computational.

> A serious autonomous system is not one that can do anything. It is one that can do a clearly defined set of things without waiting for a human, while remaining inside explicit permission, evidence, escalation, and accountability boundaries.

**SECTION TAKEAWAY**

**Main idea.** Autonomous agency begins when the system can initiate bounded workflows in response to events rather than waiting for a human prompt.

**What this teaches us about AI systems.** Autonomy is fundamentally about delegated initiative and authority. The key design variables are scope, permissions, thresholds, escalation, and reversibility.

**Main takeaway for financial modeling with AI.** Risk, liquidity, compliance, research, and market-monitoring systems can move from passive analytics to event-driven agentic architectures, but every layer of autonomy should be tied to explicit financial limits and human approval rules.

## D. Institutional Knowledge Engineering: Intelligence Needs Institutional Memory

As agents become more capable, a fundamental limitation emerges: general intelligence does not imply knowledge of the institution. A frontier model may know a great deal about accounting, portfolio theory, or corporate finance while knowing nothing about why a specific investment committee rejected a transaction last year, how the institution treats a recurring covenant exception, which internal data source is trusted, or what historical rationale sits behind an existing portfolio position.

Uber confronted this problem directly. The company describes an AI Context Graph containing roughly 24 million nodes and 80 million edges across more than 30 internal systems [1]. The graph connects services, teams, incidents, pull requests, design documents, deployments, datasets, and historical usage. The practical effect is striking: in one example, a graph-grounded agent located the correct answer in 38 seconds, while an ungrounded agent spent more than 20 minutes exploring, spawning subagents, encountering errors, and ultimately reaching the wrong conclusion [1].

The lesson is central:

$$\text{Model Capability} \neq \text{Institutional Capability}.$$

A more realistic relationship is

$$\text{Effective Institutional Intelligence} = f(\text{Model}, \text{Knowledge}, \text{Context}, \text{Tools}, \text{Memory}, \text{Governance}).$$

This argument is consistent with the broader literature on retrieval-augmented generation [4], but the institutional problem is larger than document retrieval. Retrieval asks, “Where is the relevant document?” Knowledge engineering asks, “How does this fact relate to the institution’s other facts, decisions, actors, assets, rules, and outcomes?”

A financial knowledge graph might connect

$$\text{Issuer} \leftrightarrow \text{Security} \leftrightarrow \text{Portfolio} \leftrightarrow \text{Mandate} \leftrightarrow \text{Risk},$$

while also connecting

$$\text{Issuer} \leftrightarrow \text{Research} \leftrightarrow \text{Analyst} \leftrightarrow \text{Decision} \leftrightarrow \text{Outcome}.$$

At that point, the institution begins to possess a machine-accessible memory. This is strategically important because two firms may rent the same frontier model but obtain very different results if one has encoded its institutional history, transaction precedents, analytical relationships, and decision context far more effectively.

There is also an economic implication. Poor context forces the agent to search. Search consumes tokens, time, tool calls, and attention. Better context therefore improves not only accuracy but the economics of intelligence.

**Executive Principle 3: Institutional Knowledge Is Part of the Model's Effective Intelligence**

A stronger model cannot compensate indefinitely for poor organizational memory. The institution that makes its own knowledge accessible, structured, permissioned, and traceable may extract more value from the same underlying model.

**SECTION TAKEAWAY**

**Main idea.** Institutional AI becomes powerful when general model intelligence is grounded in the organization's own memory, relationships, history, and evidence.

**What this teaches us about AI systems.** Knowledge architecture is not a peripheral database feature. It changes search cost, grounding, provenance, and the quality of reasoning.

**Main takeaway for financial modeling with AI.** The proprietary edge may increasingly come from linking models to issuer histories, research, prior forecasts, portfolio holdings, mandates, committee decisions, transactions, and realized outcomes rather than from the forecasting model alone.

## E. Cybersecurity and Harnessing: Once AI Can Act, Security Must Govern Action

The emergence of tool-using agents fundamentally changes the cybersecurity problem. When a model only produces text, the major concerns include hallucination, confidentiality, prompt injection, manipulation, and inappropriate output. Once an agent can query databases, modify files, call trading systems, communicate externally, or initiate transactions, the risk surface expands from information to action.

Uber's architecture provides a useful design principle. More than 1,000 Model Context Protocol servers are routed through a unified gateway that centralizes authentication and policy enforcement [1]. The desired architecture is therefore

$$\text{Agent} \rightarrow \text{Control Gateway} \rightarrow \text{Authorized Capability},$$

rather than

$$\text{Agent} \to \text{Unrestricted Environment.}$$

This is especially important in finance. A financial agent may interact with market data, portfolio systems, client records, payment rails, trading infrastructure, confidential research, pricing engines, legal documents, or regulatory reporting. The core security question becomes: what is the agent authorized to observe, invoke, modify, approve, and execute?

This is where the concept of the *harness* becomes central. The harness is the environment surrounding the model that determines identity, permissions, available tools, context, evaluation, observability, and rules. A simplified representation is

$$H = \{I, P, T, C, E, O, R\},$$

where $I$ is identity, $P$ permissions, $T$ tools, $C$ context, $E$ evaluation, $O$ observability, and $R$ rules.

The model may generate possible actions. The harness determines which actions are institutionally permissible. This gives us an important governance principle: in autonomous AI, control often comes less from attempting to supervise every internal thought than from governing what the system can see, call, change, approve, and execute.

Financial institutions already understand this logic. Trading mandates impose position limits. Payment systems use approval hierarchies. Operational systems enforce segregation of duties. Access is role-based. Important decisions leave audit trails. Agentic AI does not make these concepts obsolete; it extends them to machine actors.

The risk object therefore also expands. Traditional model-risk management focuses on assumptions, calibration, inputs, performance, and monitoring. Agentic system risk includes the model but also context, tools, orchestration, permission boundaries, external data, workflow logic, and interactions among agents. A competent model can still produce institutional failure if it calls the wrong tool, retrieves manipulated context, exceeds its authority, or recursively amplifies a bad decision.

The relevant governance object is no longer only the model:

$$\text{Model} + \text{Context} + \text{Tools} + \text{Agent} + \text{Workflow} + \text{Environment.}$$

High-stakes finance requires testing the full execution path.

**SECTION TAKEAWAY**

**Main idea.** Once AI can act, cybersecurity and governance must constrain actions, permissions, tools, and execution paths rather than focusing only on model outputs.

**What this teaches us about AI systems.** The harness is the operational constitution of the agent. It defines identity, access, observability, evaluation, and where machine authority ends.

**Main takeaway for financial modeling with AI.** A financial model connected to real institutional systems must operate within the same seriousness of controls applied to trading limits, payment authority, segregation of duties, provenance, and auditability.

## F. Self-Improving AI: When the System Begins Modifying Its Own Procedures

Perhaps the most consequential part of Uber's case is the move toward what it calls Continuous Skill Improvement. The idea is to capture problems encountered during skill execution and use those traces to improve the skills over time [1]. The architecture therefore changes from

$$\text{Human Designs Procedure} \rightarrow \text{Machine Executes}$$

to

$$\text{Execute} \rightarrow \text{Observe} \rightarrow \text{Evaluate} \rightarrow \text{Modify} \rightarrow \text{Execute Again.}$$

Suppose an agent uses policy $\pi_t$ and produces outcome $R_t$. An improvement mechanism may update the procedure according to

$$\pi_{t+1} = \pi_t + \Delta\pi(R_t, \tau_t),$$

where $\tau_t$ represents the execution trace. The architecture now learns not only facts but how to perform work.

This has an evolutionary character. Different procedures can be generated, compared, selected, modified, and retained. In conceptual terms, variation, evaluation, selection, and retention begin to occur within the operational system. The connection to reinforcement learning is natural [5], although production agent systems may improve through several mechanisms: trace analysis, automated evaluation, prompt revision, workflow changes, tool selection, routing changes, and skill rewriting.

For finance, the opportunity is substantial. A fraud-detection agent may learn from newly observed attacks. A credit agent may discover that certain combinations of signals are associated with poor historical outcomes. A risk agent may refine escalation rules after reviewing prior incidents. A research system may learn which evidence sequences are most useful for distinguishing signal from narrative noise.

The governance challenge grows at the same time. If the procedure itself can change, the institution must know what changed, why it changed, which evidence justified the update, how the new version was tested, and whether the change can be rolled back. Self-improvement therefore makes provenance more important, not less.

**SECTION TAKEAWAY**

**Main idea.** Self-improving systems learn not only what answer to produce but how the task itself should be performed.

**What this teaches us about AI systems.** Execution traces become a new learning asset. Skills, routing, procedures, and agent behavior can improve through repeated evaluation.

**Main takeaway for financial modeling with AI.** Financial AI systems should preserve complete execution histories and realized outcomes so that forecast errors, false alarms, missed signals, successful interventions, and human overrides can become controlled learning material.

## Game Theory, Cooperation, and Emergent Multi-Agent Behavior

Once an institution contains multiple agents that interact repeatedly, the analytical problem changes again. Agents may cooperate, compete, specialize, delegate, share information, form coalitions, or alter their behavior in response to one another. At this stage, machine learning alone is no longer sufficient to describe the system. Game theory becomes relevant.

Let the set of agents be

$$N = \{1, 2, \ldots, n\}.$$

A coalition $S \subseteq N$ produces value $v(S)$. If two agents working together create more value than they would independently,

$$v(\{i, j\}) > v(\{i\}) + v(\{j\}),$$

then cooperation creates surplus. Cooperative game theory provides a language for studying how such surplus arises and how it might be distributed [7]. Multi-agent systems research provides the broader computational context [9].

The particularly interesting question is whether collaboration must always be explicitly prescribed. As agents become more capable, persistent, and adaptive, they may discover useful divisions of labor on their own. One agent may learn that another is particularly reliable for a certain class of task. Specialist agents may begin forming recurring coalitions. Delegation patterns may stabilize. Information-sharing equilibria may emerge. The system can therefore develop organizational behavior that is not obvious from the design of any individual agent.

This is where complex-systems thinking also becomes relevant. Holland's work on adaptive systems emphasizes how macro-level organization can arise from interacting local actors [8]. Once a financial institution contains a population of agents that operate continuously, learn from outcomes, and interact strategically, the metaphor of a workflow may become too narrow. The system begins to resemble an ecology.

For finance, the research implications are unusually rich. A macro agent, liquidity agent, fundamental agent, credit agent, market-microstructure agent, and portfolio-risk agent may cooperate in ways that improve collective performance. Yet their incentives and interactions also require governance. They may over-rely on one another, amplify a shared error, create herding, or form coordination patterns that were not intended by the designers.

The final frontier is therefore not simply autonomous agents. It is autonomous *populations* of agents whose collective behavior may itself become a source of intelligence or risk.

**Executive Principle 4: Govern the Population, Not Only the Individual Agent**

As multi-agent systems become persistent and adaptive, institutional risk can arise from interaction effects even when every individual agent appears well behaved in isolation.

**SECTION TAKEAWAY**

**Main idea.** Repeated interaction among agents introduces cooperation, competition, specialization, coalition formation, and emergent behavior.

**What this teaches us about AI systems.** Multi-agent AI should be analyzed with machine learning, game theory, and complex-systems thinking because collective behavior may differ fundamentally from individual-agent behavior.

**Main takeaway for financial modeling with AI.** Future financial modeling may involve populations of specialist agents whose cooperation improves decisions, but the institution will need to study their incentives, interaction patterns, shared failure modes, and emergent equilibria.

## The Six-Layer Architecture for Financial Professionals

The Uber case can now be mapped into a six-layer architecture for understanding institutional AI.

| Layer | Central Question | Financial Interpretation |
|---|---|---|
| **A. System** | What constitutes the AI capability? | Models, context, tools, data, evaluation, orchestration, and governance. |
| **B. Agents** | How does AI perform multi-step work? | Skills, tool use, reasoning loops, delegation, and specialist agents. |
| **C. Autonomy** | Who initiates action? | Event-triggered monitoring, bounded action, thresholds, and escalation. |
| **D. Knowledge** | What does the institution know? | Institutional memory, knowledge graphs, research history, provenance, and context. |
| **E. Cybersecurity** | What may the agent do? | Identity, permissions, least privilege, tool control, logging, and auditability. |
| **F. Self-Improvement** | How does the system evolve? | Trace-driven improvement, multi-agent cooperation, game theory, and emergence. |

Table 1: A six-layer architecture for understanding institutional AI in finance.

These layers should not be treated as disconnected topics. They form a progression:

$\boxed{\text{SYSTEM}} \rightarrow \boxed{\text{AGENT}} \rightarrow \boxed{\text{AUTONOMY}} \rightarrow \boxed{\text{KNOWLEDGE}} \rightarrow \boxed{\text{HARNESS}} \rightarrow \boxed{\text{ADAPTIVE MULTI-AGEN}}$

The pedagogical advantage of this architecture is that each stage makes the next one intelligible. One cannot properly understand autonomous agency without first understanding agents. One cannot safely deploy autonomous agents without knowledge, identity, tools, and governance. Once persistent agents begin learning and interacting, game theory and complex-systems thinking become unavoidable.

This is also why AI education for financial professionals should not stop at prompting. Prompting is useful, but it occupies only the entrance to the subject. The increasingly important competencies are architectural: understanding how models are composed, how tools are invoked, how knowledge is grounded, how authority is bounded, how outcomes are evaluated, how systems learn, and how populations of agents behave.

**SECTION TAKEAWAY**

**Main idea.** The six layers form a coherent progression from AI-system literacy to adaptive, governed multi-agent institutions.

**What this teaches us about AI systems.** AI education should move from models to systems, from systems to agents, from agents to autonomy, and then toward institutional knowledge, governance, self-improvement, and strategic interaction.

**Main takeaway for financial modeling with AI.** A serious financial AI curriculum should combine statistical and reasoning models with agent architecture, knowledge engineering, optimization, cybersecurity, governance, and multi-agent cooperation.

## The Autonomous Financial Factory: A Concrete Illustration

Consider how the architecture could appear inside institutional investment management. The system continuously observes market prices, financial statements, macroeconomic data, regulatory filings, news, alternative data, portfolio exposures, and internal research. Its institutional knowledge layer understands the relationships connecting issuers, securities, portfolios, historical decisions, analysts, mandates, counterparties, and risk limits.

A meaningful event occurs: perhaps an earnings surprise, a sudden widening in credit spreads, a geopolitical shock, or a material change in liquidity. The event activates an orchestrating agent. That agent does not immediately "make the investment decision." Instead, it decomposes the problem.

A fundamental-analysis agent evaluates earnings quality, guidance, margins, and valuation. A credit agent inspects debt structure and refinancing risk. A market agent studies price formation, liquidity, and order-book behavior. A macro agent evaluates scenario exposure. A risk agent measures contribution to portfolio volatility, drawdown, factor exposure, and concentration. A knowledge agent retrieves the historical investment thesis, prior committee discussion, and previous forecast errors.

The orchestrator compares the evidence. Deterministic financial tools perform calculations that should not be left to generative reasoning. A portfolio optimizer evaluates feasible allocations subject to mandate, liquidity, concentration, and risk constraints. The system produces a recommendation with provenance.

If the recommendation lies inside a narrow pre-authorized band, a bounded action may be executed automatically. If the decision is larger, unusual, or outside policy, the system escalates. Every data source, model version, tool call, recommendation, approval, and action is logged. The eventual outcome is added to the institution's learning environment.

That architecture is no longer an AI assistant. It is the beginning of an autonomous financial operating system.

The critical point is not that every institution will adopt precisely this structure. The point is that most of the underlying components now exist independently. The next wave of financial innovation may therefore come less from inventing a completely new model than from composing existing forms of intelligence into institution-specific systems.

**SECTION TAKEAWAY**

**Main idea.** The autonomous financial factory combines continuous observation, specialist agents, institutional knowledge, deterministic modeling, constrained optimization, governed action, and learning.

**What this teaches us about AI systems.** The real power of modern AI comes from composition. Models become far more consequential when linked to data, tools, memory, controls, and other agents.

**Main takeaway for financial modeling with AI.** The frontier is not a single superior forecasting model. It is an integrated environment in which multiple forms of intelligence jointly observe, analyze, optimize, act, document, and learn under explicit financial governance.

## The Economics of the AI Factory: Architecture Determines Cost

Uber's article is also a useful study in the economics of machine cognition. The company decomposes AI spending into scale, efficiency, and unit-cost drivers and emphasizes that rapid usage growth does not have to imply proportionate cost growth when architecture improves [1]. This matters because an AI workflow can multiply calls invisibly. A single user request may launch several agents, each of which invokes models, tools, retrieval, and additional subagents.

The visible price of the model is therefore not the economic unit the institution should optimize.

For finance, the correct denominator is the business outcome. Instead of asking only about dollars per million tokens, management should ask about

$$\frac{\text{AI Cost}}{\text{Validated Investment Analysis}}, \quad \frac{\text{AI Cost}}{\text{Credit Review}}, \quad \frac{\text{AI Cost}}{\text{Risk Case Resolved}},$$

or any other unit corresponding to a useful, controlled financial result.

This perspective also reveals why architecture is an economic decision. Better retrieval can reduce expensive context. Smaller specialist models can handle routine extraction. Deterministic tools can perform calculations more reliably and cheaply than a frontier reasoning model. A stronger model can be reserved for the part of the workflow where sophisticated judgment has the highest marginal value. Human review can be concentrated on exceptions rather than repeated across every case.

The result is a form of computational capital allocation. Expensive intelligence should be allocated where it produces the greatest incremental value.

This is familiar territory for financial professionals. The challenge is to apply the same discipline used in capital allocation to machine cognition itself.

**Executive Principle 5: Optimize Cost per Useful, Controlled Outcome**

Token price is a component of cost, not the business case. The architecture determines how much expensive intelligence is consumed per completed financial task.

**SECTION TAKEAWAY**

**Main idea.** The economics of AI are determined by the complete workflow, not by the headline price of the model.

**What this teaches us about AI systems.** Routing, retrieval, caching, specialist models, deterministic tools, agent depth, and human escalation determine the marginal cost of machine cognition.

**Main takeaway for financial modeling with AI.** Financial institutions should measure the cost per validated analysis, risk intervention, credit decision, research output, or transaction workflow and optimize architecture around that unit.

## Architecture as Competitive Advantage

One of the most important strategic implications follows from the increasing accessibility of frontier models. If many banks, asset managers, exchanges, and investment firms can rent access to similar general-purpose intelligence, then the model itself cannot remain the sole source of sustainable differentiation.

The competitive asset increasingly moves into the architecture surrounding the model. Institutional knowledge can differ. Proprietary skills can differ. Internal data can differ. Agent orchestration can differ. Evaluation can differ. Permission structures can differ. The institution's ability to learn from its own execution traces can differ.

Institutional AI capability can therefore be represented conceptually as

$$\mathcal{I} = f(M, K, S, A, H, G, L),$$

where $M$ is model capability, $K$ institutional knowledge, $S$ skills, $A$ agent architecture, $H$ harness quality, $G$ governance, and $L$ learning capacity.

As model access becomes more interchangeable, the relative importance of the other terms may rise.

This is particularly important in finance because financial institutions have always competed through information advantage, analytical process, execution, judgment, distribution, balance sheet, brand, and institutional memory. Agentic architecture may become another form of organizational capital. A firm that has deeply encoded its investment process, captured its historical decisions, built reusable financial skills, created a strong evaluation environment, and designed robust agent governance may extract far more value from the same model available to competitors.

Owning the model is therefore not the same as owning the advantage.

**SECTION TAKEAWAY**

**Main idea.** As frontier models commoditize, competitive advantage is likely to migrate toward proprietary knowledge, workflows, skills, controls, and learning architecture.

**What this teaches us about AI systems.** Institutional capability is a function of the entire architecture, not only the intelligence engine.

**Main takeaway for financial modeling with AI.** Financial firms should treat knowledge architecture, reusable agent skills, evaluation frameworks, provenance, and governance as strategic assets rather than as implementation details.

## What Financial Professionals Need to Learn

The educational implication is straightforward but demanding. AI education for financial professionals cannot stop at prompt engineering. Prompting is useful because it allows professionals to access model capabilities quickly. But precisely because the interface has become easier, the architecture behind the interface has become more important to understand.

The financial professional increasingly needs to know how an agent decides what to do next, how it discovers tools, how it retrieves institutional context, how it coordinates with specialized agents, how authority is bounded, how actions are traced, how performance is evaluated, and how procedures can improve over time. Once multiple agents coexist, the professional also needs to understand strategic interaction, collective failure, and emergent behavior.

The required intellectual toolkit therefore expands beyond traditional machine learning and LLM literacy. It includes systems architecture, knowledge engineering, optimization, cybersecurity, governance, multi-agent systems, game theory, and complex adaptive systems.

This does not mean that every financial professional must become a computer scientist. The objective is architectural fluency, not engineering mastery. A CEO does not need to implement an embedding index in order to understand why institutional knowledge matters. A portfolio manager does not need to write an agent framework in order to understand why an autonomous system must have execution limits. A board member does not need to understand transformer mathematics in order to ask whether critical actions are reconstructable, attributable, and reversible.

The paradox is that the easier AI becomes to use, the more important it becomes to understand what lies behind the interface.

**SECTION TAKEAWAY**

**Main idea.** Ease of use should not be confused with ease of understanding. The interface simplifies while the institutional architecture becomes richer.

**What this teaches us about AI systems.** The next generation of AI literacy is architectural literacy: knowing how models, tools, agents, knowledge, evaluation, and governance combine.

**Main takeaway for financial modeling with AI.** Financial professionals need enough system literacy to design, supervise, challenge, and govern AI-enabled financial models and workflows without needing to become machine-learning engineers.

## Conclusion: From Software Factory to Cognitive Institution

Uber's Software Factory should be interpreted as more than an experiment in engineering productivity. It is an early empirical description of an organizational form in which artificial intelligence is becoming part of the operating fabric of the institution.

Agents proliferate. Skills become reusable. Orchestrators divide work among subagents. Institutional knowledge is engineered explicitly. Tool access is placed behind policy gateways. Managed agents initiate work. Evaluation becomes benchmark-driven. Costs are measured at the level of completed outcomes. Execution traces begin feeding back into skill improvement. Each development is important individually, but together they imply that the institution itself is changing.

For decades, organizations used software as an instrument. Humans remained the principal cognitive actors. Software stored information, executed calculations, and automated predefined procedures. Agentic AI begins to blur that boundary. The institution increasingly contains computational actors capable of interpreting context, pursuing objectives, selecting tools, delegating work, evaluating outcomes, and initiating bounded actions.

Institutional knowledge becomes machine-accessible. Governance becomes architectural. Cybersecurity becomes inseparable from autonomy. Learning begins to modify procedures. Interaction among agents introduces cooperation, competition, coalition formation, and emergence.

The progression can therefore be summarized as

$$\text{Software} \rightarrow \text{Intelligent Software} \rightarrow \text{Agents} \rightarrow \text{Autonomous Agents} \rightarrow \text{Knowledge-Grounded Institution} \rightarrow \text{Adapt}$$

The deepest lesson for finance is not that every professional will receive a better chatbot. It is that artificial intelligence may alter the architecture through which financial work itself is organized.

The future bank, investment manager, exchange, insurer, investment bank, or regulator may increasingly become a hybrid cognitive system in which humans, models, institutional knowledge, deterministic analytics, optimization, software, and autonomous agents operate together. In that environment, the most valuable organization may not be the one with the strongest model. It may be the one that has learned how to *harness intelligence as an institution*.

**The final financial rule**

## Do not build the financial institution of the AI age around a chatbot.

Build it around a governed cognitive architecture:
models, knowledge, skills, agents, controls, learning, and human judgment.

**SECTION TAKEAWAY**

**Main idea.** The next stage of AI is institutional and architectural. Machine intelligence is becoming embedded in the operating structure of organizations.

**What this teaches us about AI systems.** Models, agents, tools, knowledge, governance, feedback, and cooperation are converging into persistent cognitive systems.

**Main takeaway for financial modeling with AI.** The long-term opportunity is to build financial architectures in which models no longer operate as isolated analytical engines but as governed components of intelligent, knowledge-rich, adaptive institutions.

## References

## References

[1] U. K. Medisetty, "Running a Software Factory Efficiently at Uber Scale," *Uber Engineering*, August 2026.

[2] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, "ReAct: Synergizing Reasoning and Acting in Language Models," *International Conference on Learning Representations*, 2023.

[3] Q. Wu et al., "AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation," *arXiv:2308.08155*, 2023.

[4] P. Lewis et al., "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," *Advances in Neural Information Processing Systems*, vol. 33, 2020.

[5] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed., MIT Press, 2018.

[6] D. Amodei, C. Olah, J. Steinhardt, P. Christiano, J. Schulman, and D. Mané, "Concrete Problems in AI Safety," *arXiv:1606.06565*, 2016.

[7] M. J. Osborne and A. Rubinstein, *A Course in Game Theory*, MIT Press, 1994.

[8] J. H. Holland, *Adaptation in Natural and Artificial Systems*, MIT Press, 1992.

[9] M. Wooldridge, *An Introduction to MultiAgent Systems*, 2nd ed., Wiley, 2009.

[10] M. Zaharia et al., "The Shift from Models to Compound AI Systems," *Berkeley Artificial Intelligence Research*, 2024.